



# Why AMI is needed, what are the advantages, where have you used it ?

## What is an AMI (Amazon Machine Image)?

An **Amazon Machine Image (AMI)** is a **pre-configured template** that contains everything needed to **launch an EC2 instance**, including:

- The **OS** (e.g., Amazon Linux, Ubuntu, Windows)
  - Application **server**, libraries, and dependencies
  - Pre-installed **software** or custom configurations
  - Any **user-defined settings** or startup scripts
- 

## Why is AMI Needed?

Imagine you're setting up a server with:







- Ubuntu 22.04
- Apache web server
- Python 3.11
- A Django application
- Custom firewall rules

Once you've done all that, you don't want to repeat the setup every time you launch a server. So, you **create an AMI** of this configured instance.

Now, every time you launch an EC2 from this AMI, **everything is already set up** — OS, software, configurations — in minutes.

---

## Advantages of AMI

Advantage	Explanation
 <b>Reusability</b>	Quickly launch multiple identical instances using the same AMI.
 <b>Fast Scaling</b>	Easily spin up pre-configured servers for auto-scaling during traffic spikes.
 <b>Backup &amp; Recovery</b>	Take AMI snapshots as backups of your EC2 instance — useful before major changes.
 <b>Custom Configuration</b>	Tailor the image with your own app, tools, and security policies.
 <b>Cross-region Deployment</b>	Copy AMIs across regions for global deployment.
 <b>Team Standardization</b>	Developers and testers use the exact same environment, reducing "it works on my machine" issues.

---

## Example Use Case

**Scenario:** An e-commerce company sets up a web server with:

- NGINX + Node.js + MongoDB
- SSL certificate and monitoring tools

Instead of setting this up every time, they **create an AMI** called `Ecom-WebServer-2025`.

When auto-scaling triggers (e.g., during a festive sale), new instances **launch instantly from this AMI**, serving users in seconds.

## Interview Tip

“I usually create custom AMIs after setting up base environments or before deploying major updates, so that I can roll back easily or replicate the setup across staging and production.”

---

Mistakes/Lessons learnt:

*1. Earlier, I used to create AMIs without proper tagging. Later, it became difficult to track which AMI belonged to which environment or release.*

2. Initially, I didn't delete old unused AMIs and snapshots, which led to unnecessary storage costs. Now, I automate AMI cleanup with lifecycle policies.
3. Once I forgot to update the AMI after a critical patch, and newly launched instances were missing the fix. Since then, I've made **AMI refresh** part of our **deployment checklist**.

#### A sample AMI creation command

```
aws ec2 create-image \  
  --instance-id i-0abc1234def5678gh \  
  --name "MyApp-Image-April2025" \  
  --description "An AMI for my app as of April 2025" \  
  --no-reboot
```

#### How do you create a custom AMI and when would you use one?

Use `create-image` command or console; used for consistent deployments, backups, or launching multiple identical servers.

#### 1. What are the differences between public, private, and shared AMIs?

- **Public:** Available to all AWS users.
- **Private:** Visible only to the owner.
- **Shared:** Explicitly shared with specific AWS accounts.

#### 2. Can you create an AMI without stopping your instance?

Yes, using `--no-reboot` option, but the AMI might miss in-memory data or file system changes.

##### 1. In-Memory Data

- **Application cache** (e.g., Redis, Memcached running on RAM)
- **Session data** stored in memory by application servers (like Node.js, Python, Java apps)
- **Running database transactions** in memory (e.g., data in a PostgreSQL, MySQL buffer pool not yet flushed to disk)
- **Log files** that are buffered in memory and not yet flushed to disk
- **Temporary processing data** held in RAM (e.g., image processing jobs, video transcoding data)

---

##### 2. File System Changes (Not Flushed Yet)

- Files that were created or modified but still sitting in **OS disk cache** and not flushed to disk (especially if `fsync` wasn't called)

- **Database writes** that are still in the buffer/cache (example: uncommitted writes in MySQL's InnoDB buffer pool)
- **Active downloads/uploads** where parts of the file are not yet written to disk
- **Ongoing software installations/updates** that have not yet finalized disk writes

3. **What happens when you launch an EC2 instance from a custom AMI?**

A new EC2 instance is created with the exact OS, configuration, and data as in the AMI.

4. **How do you ensure an AMI is shared securely between accounts?**

Share it using `modify-image-attribute` and also share the snapshot; avoid making it public.

5. **How would you automate AMI creation and retention policies?**

Use AWS Lambda with CloudWatch Events or AWS Systems Manager Automation with lifecycle rules.

**Scenario** Why automate AMI creation + retention?

**Nightly backup of production servers** So you always have a fresh backup AMI in case of crash

**Before critical deployments** (e.g., app updates) So you can roll back if the update fails

**For compliance and audits** Certain industries require regular snapshots and retention for X days

6. **What limitations or considerations apply when copying AMIs across regions?**

You must copy the AMI manually; ensure associated snapshots are available in the destination region.

An **AMI (Amazon Machine Image)** created in one AWS region is **specific to that region** — it is **not automatically available** in other regions.

```
aws ec2 copy-image --source-region us-east-1 --source-image-id ami-1234567890abcdef0 --name "MyCopiedAMI" --region us-west-2
```

**Important Points:**

- The snapshot linked to the AMI will also be copied automatically.
- Copying across regions may incur additional costs (for snapshot storage and transfer).
- Make sure that any associated permissions (e.g., if AMI is shared) are also handled after copying.

7. **How do AMIs work in Auto Scaling Groups or Launch Templates?**


The AMI ID is referenced in the launch template/configuration to spin up identical instances automatically.

8. **Can you update an existing AMI? If not, what's the best practice?**

No, AMIs are immutable. Best practice is to launch from the old AMI, make changes, and create a new AMI.

1. **What does an Amazon Machine Image (AMI) include?**

- A. Only the operating system
- B. The OS, application server, and application data
- C. Only the instance type
- D. Only the snapshot of a volume

 **Answer:** B. The OS, application server, and application data

---

2. **Which AWS CLI command is used to create an AMI?**

- A. `aws ec2 start-image`
- B. `aws ec2 build-image`
- C. `aws ec2 create-image`
- D. `aws ami create`

 **Answer:** C. `aws ec2 create-image`