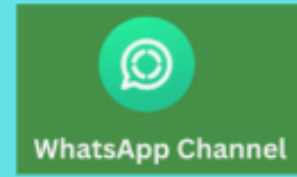


@devopschallengehub



Interview questions on Advanced IAM Features

What are permission boundaries in IAM?

https://docs.aws.amazon.com/IAM/latest/UserGuide/access_policies_boundaries.html

Permission boundaries in AWS Identity and Access Management (IAM) are **advanced features** that set the **maximum permissions** an IAM role or user can have. Even if the IAM user or role is granted more permissions through their policies, **they cannot exceed what's defined in the permission boundary**.

Think of it like this:

- **IAM Policy** = Grants permissions
- **Permission Boundary** = Restricts permissions to a maximum limit

So, **effective permissions** = **intersection of IAM policy and permission boundary**

✅ Use Case:

Imagine you're allowing developers to create their own IAM roles via automation. You want to ensure that they **can't give themselves full admin rights**, even if they try. You apply a **permission boundary** that says:

“You can create roles, but those roles can only manage S3 and Lambda, not EC2 or IAM.”

📌 Example:

Suppose a user has this **IAM policy**:

```
json
-----
{
  "Effect": "Allow",
  "Action": "s3:*",
  "Resource": "*"
}
```

And the **permission boundary** is:

```
json
-----
{
  "Effect": "Allow",
  "Action": [
    "s3:GetObject",
    "s3:ListBucket"
  ],
  "Resource": "*"
}
```

Then, even though the IAM policy allows all S3 actions, **the user will only be allowed to perform** `GetObject` **and** `ListBucket` due to the permission boundary.

Summary:

- **Permission boundary** = **upper limit** of what a user/role can do.
- They are attached to **IAM roles or users**, **not groups**.
- They're mainly used to **control delegated access** and enforce **security guardrails**.

How do IAM roles work with cross-account access?

- **Account A** has an S3 bucket).
- **Account B** wants to use it safely.
- **IAM Role** is the **gatekeeper** that lets them in — **with rules**.

Examples

◆ Step 1: Create Role in Account A

In Account A, create a role and define a **trust policy** to allow Account B to assume it.

✅ Trust Policy (Account A's role):

```
json
-----
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Principal": {
        "AWS": "arn:aws:iam::ACCOUNT_B_ID:root"
      },
      "Action": "sts:AssumeRole"
    }
  ]
}
```

This means: “I trust all users in Account B to assume this role.”

STS (Security Token Service)

◆ Step 2: Attach Permissions to the Role in Account A

Add a permissions policy to the role. For example, give access to an S3 bucket:

✅ Permissions Policy:

```
json
-----
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": ["s3:GetObject"],
      "Resource": "arn:aws:s3::my-bucket-name/*"
    }
  ]
}
```

This means: “If someone assumes this role, they can read files from my S3 bucket.”

◆ Step 3: Assume the Role from Account B (Using AWS CLI)

From Account B, assume the role using `aws sts assume-role`. Replace ARNs with actual values.

✅ CLI Command (from Account B):

```
bash
-----
aws sts assume-role \
  --role-arn arn:aws:iam::ACCOUNT_A_ID:role/RoleName \
  --role-session-name MySession
```

✅ Response:

You'll get temporary credentials:

```
json
-----
{
  "Credentials": {
    "AccessKeyId": "ASIAxxxxxxxxxxxx",
    "SecretAccessKey": "xxxxxxxxxxxxxxxx",
    "SessionToken": "xxxxxxxxxxxxxxxx",
    "Expiration": "2025-05-22T08:20:00Z"
  }
}
```

◆ Step 4: Use Temporary Credentials

Use those temporary keys to **access Account A's resources**.

✅ Example: Use **aws s3** with temporary keys

```
bash
-----
export AWS_ACCESS_KEY_ID=ASIAxxxxxxxxxxxx
export AWS_SECRET_ACCESS_KEY=xxxxxxxxxxxxxxxx
export AWS_SESSION_TOKEN=xxxxxxxxxxxxxxxx

aws s3 ls s3://my-bucket-name
```

This shows the files in the S3 bucket **from Account B using the role in Account A**.

Final Recap

Step	What You Do	Code
✅ Create Role in Account A	Trust Account B	JSON trust policy
✅ Add Permissions	Grant limited access	S3 read-only example
✅ Assume Role	Use <code>sts assume-role</code>	AWS CLI
✅ Use Role	Export temporary keys	Access S3

Explain how STS (Security Token Service) works for temporary credential management ? 🤖
How does AssumeRole work and when would you use it in a DevOps pipeline? Answer in bullet points, short sentences.

✅ What is STS (Security Token Service)?

- STS gives **temporary credentials** (Access Key, Secret Key, Session Token).
- Used for **short-term, limited access** to AWS resources.

- Helps avoid hardcoding permanent credentials.
 - Credentials usually last **15 mins to a few hours**.
 - Used for **cross-account access**, CI/CD pipelines, federated users, etc.
-

✓ How AssumeRole Works

- A user or app **calls** AssumeRole using the AWS CLI or SDK.
 - STS checks if the **trust policy allows** the caller.
 - If allowed, STS returns **temporary credentials**.
 - These credentials are then used to **access AWS resources**.
 - Permissions are based on the **role's permission policy**.
-

✓ When to Use AssumeRole in DevOps Pipelines

- 🔄 To **switch roles** and deploy to another AWS account (e.g., staging → prod).
- 🗑️ To ensure **least privilege** – only give access during build/deploy steps.
- 🔑 To avoid **storing long-term credentials** in CI/CD tools (like Jenkins, GitHub Actions).
- 📦 To allow pipelines to access **specific services** (e.g., push to ECR, deploy to ECS).
- 📁 To fetch or store files in **S3 buckets across accounts** during automation.
- 🧪 To test apps in **isolated environments** (each with different roles).

What are service control policies (SCPs)? When are they used?

✓ What are Service Control Policies (SCPs)?

- SCPs are part of **AWS Organizations**.
- They set **permission guardrails** for AWS accounts.
- SCPs **do not grant** permissions — they **only restrict** what can be done.
- Think of them as **"master switch" rules** across accounts.
- They apply to **all IAM users, groups, and roles** in the account.

✅ When Are SCPs Used?

- 🔒 To **enforce security boundaries** across multiple accounts.
 - 🚫 To **block risky services** (e.g., prevent use of EC2, IAM in dev accounts).
 - ✅ To allow only specific regions (e.g., "Only use ap-south-1").
 - 📦 To limit services based on **environment** (e.g., no RDS in test account).
 - 🛡️ To ensure **compliance** and **centralized control** in large organizations.
 - 🧑‍💻 To prevent account admins from overriding **org-wide policies**.
-

Even if an IAM user has full AdministratorAccess, **SCP can still block** that action.

What are permission boundaries in IAM?

Q1: What is the main purpose of a permission boundary in IAM?

- A. To give full admin access to an IAM user
- B. To define the maximum permissions an IAM role or user can get
- C. To restrict access to AWS billing
- D. To block access to specific IP addresses

✅ **Correct Answer: B**

How do IAM roles work with cross-account access?

Q2: What allows an IAM role in Account A to be used by a user in Account B?

- A. IAM policy in Account B
- B. CloudTrail logging
- C. Trust policy in Account A's role
- D. Billing dashboard

✅ **Correct Answer: C**

Which of the following is a valid use case for AssumeRole in DevOps?

- A. Monitoring CPU usage
- B. Deploying resources to another AWS account securely**
- C. Managing CloudFront cache
- D. Encrypting EBS volumes

 **Correct Answer: B**

What are Service Control Policies (SCPs)? When are they used?

Q6: What is true about SCPs in AWS Organizations?

- A. SCPs give users full access to all services
- B. SCPs set permission guardrails across accounts
- C. SCPs are for billing settings only
- D. SCPs are IAM policies applied to S3 buckets

 **Correct Answer: B**

What happens if an SCP blocks an action but an IAM policy allows it?

- A. The action is allowed
- B. The IAM policy wins
- C. The action is denied**
- D. AWS ignores both policies

 **Correct Answer: C**