



Question 10: How do you handle security and access control in CodePipeline?

Security is non-negotiable in any CI/CD pipeline, especially in enterprise environments where multiple teams, accounts, and environments are involved. In my DevOps implementations, I place strong focus on **IAM roles**, **least privilege**, and **secure access delegation** across accounts.

1. IAM Roles in CodePipeline: Types & Use Cases

🔧 a) Service Roles (Execution Roles)

- These are the roles **assumed by CodePipeline** or **CodeBuild** when performing actions like:
 - Storing artifacts in S3
 - Deploying using CodeDeploy
- We use **fine-grained permissions** to follow **least privilege principle**.

📌 Example: CodePipeline role can access only:

```
json
"Action": [
  "s3:GetObject",
  "codebuild:StartBuild",
  "codedeploy:CreateDeployment"
],
"Resource": [
  "arn:aws:s3:::my-app-artifacts/*",
  "arn:aws:codebuild:region:account:project/my-project"
]
```

b) Cross-Account Roles

- When we had **multiple AWS accounts** (e.g., dev, prod, security), we used **cross-account IAM roles** so pipelines in one account can trigger deployments in another securely.

Real Use Case:

Pipeline in **Account A** needed to deploy to ECS in **Account B**.

So we created a **trusted role in Account B** that allowed CodePipeline from Account A to assume it.

Trust Policy in Account B:

json

```
"Statement": [
  {
    "Effect": "Allow",
    "Principal": {
      "AWS": "arn:aws:iam::ACCOUNT-A:role/CodePipelineRole"
    },
    "Action": "sts:AssumeRole"
  }
]
```

In CodePipeline (Account A), we used assume-role action to take on this temporary role for deploying.

c) Developer IAM Access Roles

We restricted developers using:

- **View-only IAM policy** (can view pipelines, logs, execution history)
- **Approval-only role** (can only approve/reject manual approval actions)



This helped us implement **Separation of Duties (SoD)**:

- Dev can't directly deploy to production
- Only release managers can approve

2. Additional Security Best Practices

Best Practice	Description
✅ S3 Encryption	Artifacts stored in S3 are encrypted using either SSE-S3 or SSE-KMS
✅ VPC Endpoints	CodeBuild/CodeDeploy access resources securely over private AWS network
✅ Secrets Manager	API tokens, GitHub credentials, DB secrets stored securely and rotated
✅ CloudTrail Logs	Every IAM role assumption and pipeline action is logged for audit
✅ SSM Parameter Store	Used for non-secret config values (e.g., ENV name, region)



Example: GitHub Token Handling

We needed to authenticate CodePipeline with GitHub. Instead of hardcoding credentials, we:

- Stored token in **AWS Secrets Manager**
- Used **Secrets Manager integration in CodeBuild**
- Set up an IAM policy that only allowed **read access to that specific secret**

Takeaway

By using scoped IAM roles, cross-account delegation, and best practices like encryption, secrets management, and VPC isolation, we made our pipelines **secure, auditable, and compliant** — without slowing down development."

Question 11: How do you optimize CodePipeline performance and cost?

As our team scaled CI/CD across multiple microservices, we saw our build times increasing and monthly costs creeping up. That's when we did a full audit and implemented several strategies to **boost performance and reduce cost**, without compromising on quality or security.

Performance Optimization Strategies

1 Enable Caching in CodeBuild

Caching significantly reduces build times by avoiding redundant work.

Type	Benefit
Docker Layer Caching	Speeds up builds by reusing unchanged layers – 60% faster builds
NPM or Yarn Caching	Avoids re-downloading packages every time – ~70% improvement
Maven/Gradle Caching	Reduces time for dependency resolution – 50% gain

✓ Example: Enabling NPM cache in buildspec.yml

```
yaml
cache:
  paths:
    - '/root/.npm/**/*'
```

For Docker image builds:

- Use `--cache-from` strategy with ECR
- Or use **CodeBuild Docker layer cache** (via `privileged: true` and local cache settings)

2 Parallel Execution

We split the pipeline into parallel branches for independent modules like:

- Frontend
- Backend
- Unit Tests
- Integration Tests

🔍 Result:

- Reduced pipeline time from **45 minutes to 18 minutes**

📌 Pro Tip: In CodePipeline, use **multiple CodeBuild actions in parallel** under one stage.

🔧 Cost Optimization Strategies

1 Right-size CodeBuild Instance Types

We audited our CodeBuild jobs and selected instance types based on build workload:

Instance Type	Specs	Use Case
build.general1.small	2 vCPU, 3 GB RAM	Lightweight builds, unit tests
build.general1.medium	4 vCPU, 7 GB RAM	Docker builds, integration tests
build.general1.large/2XL	8+ vCPU	ML/long-running builds only if justified

📌 You're billed per second, so even saving **30 seconds per build** adds up when you have 50+ daily builds.

2 S3 Lifecycle Policies for Artifacts

Artifacts stored in S3 can cost a lot if not cleaned up.

Strategy:

- Set **S3 lifecycle rule**:
 - Move to **infrequent access (IA)** after 7 days
 - **Delete** after 30 days

Example Rule:

```
json
{
  "ID": "ArtifactLifecycle",
  "Prefix": "artifacts/",
  "Status": "Enabled",
  "Transitions": [{
    "Days": 7,
    "StorageClass": "STANDARD_IA"
  }],
  "Expiration": {
    "Days": 30
  }
}
```

📊 Monitoring & Benchmarking

1 CloudWatch Metrics

Tracked:

- **Average build duration**
- **Cache hit/miss ratio**
- **Stage failure rates**

2 Cost Alerts (AWS Budgets)

Set:

- Daily and monthly budgets
- Alert thresholds (e.g., notify if CodeBuild > ₹500/day)

3 Custom Benchmark Dashboard

We used **CloudWatch Dashboards + Lambda + DynamoDB** to:

- Track performance improvement per microservice
- Identify build jobs with abnormal duration or failure spikes

✓ Summary of Optimizations

Area	Optimization	Result
Build Time	Caching, parallel stages	2x–3x faster
Cost	Smaller instance types	~30% monthly savings
Storage	S3 lifecycle	60–70% reduction in artifact costs
Observability	Dashboards & alerts	Proactive cost & time management

🧠 Takeaway

"CI/CD performance isn't just about speed—it's about **smart execution**. With caching, parallelism, right-sized infra, and tight monitoring, we were able to cut **build time by 60%** and save **thousands per month**, while improving release frequency and confidence."

What financial control was implemented to alert the team about budget overruns?

- A. EC2 Reserved Instances
- B. IAM policy triggers
- C. AWS Budgets with alert thresholds
- D. CloudTrail auditing

✓ **Correct Answer: C**