



How do you share resources between multiple CloudFormation stacks?

Or

What are nested stacks and cross-stack references? When would you use them?

I've used both **nested stacks** and **cross-stack references** extensively to manage scalable, modular, and reusable infrastructure using AWS CloudFormation. These approaches have been crucial in keeping our infrastructure organized and maintainable, especially in larger projects.

◆ What are Nested Stacks?

Definition: A **nested stack** is essentially a child stack embedded within a parent CloudFormation stack.

It helps break down large, complex templates into smaller, focused units.

Why I Use Them:

- To stay under the **51,200-character limit** of a template
- To **reuse infrastructure patterns** like VPC, subnets, security groups
- To allow **teams to independently own and manage** parts of the infrastructure
- To simplify **lifecycle management** (e.g., compute layer changes frequently, but the networking layer is stable)

Real Usage:

In one project, our parent template included nested stacks for:

- Networking (networking.yaml)
- Compute resources (compute.yaml)
- Monitoring setup (monitoring.yaml)

Each nested template lived in S3 and was version-controlled independently.

◆ What are Cross-Stack References?

Definition: Cross-stack references allow a stack to **export outputs**, and other stacks can **import those values**.

Why I Use Them:

- To **share common infrastructure** like VPCs, security groups, ALBs, or RDS endpoints across multiple stacks
- To support **microservices architecture**, where each service stack is loosely coupled
- To separate responsibilities — for example, networking managed by one team, applications by another



How I Share Resources Between Stacks

✓ Method 1: Cross-Stack Exports/Imports

Step 1: Export in source stack:

yaml

CopyEdit

Outputs:

VpcId:

Value: !Ref MyVPC

Export:

Name: 'networking-VpcId'

Step 2: Import in dependent stack:

yaml

CopyEdit

Resources:

MyEC2Instance:

Properties:

VpcId: !ImportValue 'networking-VpcId'

✓ Method 2: SSM Parameter Store

I also store resource info like DB endpoints in SSM and reference them using `{{resolve:ssm:/myapp/database}}`. This makes it more dynamic and loosely coupled.

✓ Method 3: Pass as Parameters

When creating Stack B, I fetch outputs from Stack A and pass them using `--parameters`.



Real-World Example: E-commerce Platform

We split the infra into:

- **Foundation Stack** (nested): VPC, subnets, shared ALB
- **Database Stack**: RDS, ElastiCache, exported endpoints
- **Application Stacks**: User/Product/Order services importing shared infra

Each stack was independently deployed and owned by a specific team.



When to Use Each Approach

Use Case	Use Nested Stacks	Use Cross-Stack References
Modularity and maintainability	✓	✓
Large templates approaching limits	✓	✗
Reusable components	✓ (standard networking, monitoring)	✓ (shared VPC, SG, DB endpoint)
Microservices architecture	✗	✓
Loose coupling between teams	✓	✓

Use Case	Use Nested Stacks	Use Cross-Stack References
Lifecycle separation	✓	✓

🧠 Best Practices I Follow

For Nested Stacks:

- Single responsibility per nested stack
- Version templates in S3
- Use DependsOn to manage stack order
- Test nested stacks individually before integration

For Cross-Stack References:

- Use descriptive export names (prefix with stack name)
- Keep exports minimal — only what's needed
- Carefully plan deletion order (consumers must be deleted first)
- Avoid circular dependencies

⚠️ Common Pitfalls

- **Changing export names:** breaks downstream stacks
- **Too many cross-stack references:** leads to tight coupling
- **Circular dependencies:** difficult to untangle
- **Forgetting S3 URL for nested stacks:** they must live in S3

🚫 When NOT to Use

- Avoid nested stacks for small or simple templates
- Avoid cross-stack references when:
 - Resources change frequently together
 - You need fully atomic deployments (better handled in one stack)

✓ Conclusion

Using **nested stacks** and **cross-stack references** has helped me build clean, maintainable infrastructure across multiple environments and teams. The key is to **balance modularity with simplicity** and use these tools when they add clarity, not complexity.

NESTED STACK

1. Child Stack (e.g., `webserver-stack.yaml`)

This stack defines the web server resources.

YAML

`webserver-stack.yaml`

AWSTemplateFormatVersion: '2010-09-09'

Description: A simple web server stack

Parameters:

VpcId:

Type: String

Description: The ID of the VPC where the web server will be deployed.

SubnetId:

Type: String

Description: The ID of the subnet where the web server will be deployed.

SecurityGroupId:

Type: String

Description: The ID of the security group to associate with the web server.

Resources:

WebServerInstance:

Type: AWS::EC2::Instance

Properties:

ImageId: ami-0abcdef1234567890 # Replace with a valid AMI ID for your region

InstanceType: t2.micro

NetworkInterfaces:

- DeviceIndex: "0"

AssociatePublicIpAddress: "true"

SubnetId: !Ref SubnetId

GroupSet:

- !Ref SecurityGroupId

Tags:

- Key: Name

Value: MyWebServer

Outputs:

WebServerPublicIp:

Description: The public IP address of the web server

Value: !GetAtt WebServerInstance.PublicIp

2. Parent Stack (e.g., parent-network-and-app-stack.yaml)

This stack defines the network and then deploys the webserver-stack as a nested stack.

YAML

parent-network-and-app-stack.yaml

AWSTemplateFormatVersion: '2010-09-09'

Description: Parent stack for network and web application deployment

Resources:

VPC:

Type: AWS::EC2::VPC

Properties:

CidrBlock: 10.0.0.0/16

EnableDnsSupport: true

EnableDnsHostnames: true

Tags:

- Key: Name

Value: MyVPC

PublicSubnet:

Type: AWS::EC2::Subnet

Properties:

VpcId: !Ref VPC

CidrBlock: 10.0.1.0/24

MapPublicIpOnLaunch: true

Tags:

- Key: Name

Value: MyPublicSubnet

WebServerSecurityGroup:
Type: AWS::EC2::SecurityGroup
Properties:
 GroupName: WebServerSG
 GroupDescription: Enable HTTP access
 VpcId: !Ref VPC
 SecurityGroupIngress:
 - IpProtocol: tcp
 FromPort: 80
 ToPort: 80
 CidrIp: 0.0.0.0/0
 Tags:
 - Key: Name
 Value: WebServerSecurityGroup

Nested Stack for the Web Server

WebServerStack:
Type: AWS::CloudFormation::Stack
Properties:
 TemplateURL: webserver-stack.yaml # Assuming both files are in the same S3 bucket or local path
 Parameters:
 VpcId: !Ref VPC
 SubnetId: !Ref PublicSubnet
 SecurityGroupId: !Ref WebServerSecurityGroup

Outputs:
WebServerErrorCheck:
 Description: Public IP of the web server from the nested stack
 Value: !GetAtt WebServerStack.Outputs.WebServerPublicIp

Explanation:

- **webserver-stack.yaml (Child Stack):** This is a self-contained template that deploys an EC2 instance. It takes VpcId, SubnetId, and SecurityGroupId as parameters, which will be provided by the parent stack. It outputs the public IP of the web server.
- **parent-network-and-app-stack.yaml (Parent Stack):**
 - It first creates a VPC, a public subnet, and a security group.
 - The crucial part is the WebServerStack resource. Its Type is AWS::CloudFormation::Stack, which indicates it's a nested stack.
 - TemplateURL: This specifies the S3 URL or local path to the child stack template (webserver-stack.yaml).
 - Parameters: This is how the parent stack passes values to the child stack. It's mapping the !Ref (references) to its own created resources (VPC, Subnet, Security Group) to the parameters expected by the webserver-stack.yaml.
 - Outputs: The parent stack can also retrieve outputs from the nested stack using !GetAtt WebServerStack.Outputs.WebServerPublicIp.

How to deploy (Conceptual):

1. Upload webserver-stack.yaml and parent-network-and-app-stack.yaml to an S3 bucket (or deploy locally if using the AWS CLI and have the files in the same directory).

2. Deploy the parent-network-and-app-stack.yaml using the CloudFormation console or AWS CLI. CloudFormation will then automatically create the nested webserver-stack as part of its deployment.

CROSS REFERENCE STACK

1. Network Stack (network-stack.yaml) - The Exporter

This stack creates the network resources and **exports** the PublicSubnetId and VpcId.

YAML

```
# network-stack.yaml
```

```
AWSTemplateFormatVersion: '2010-09-09'
```

```
Description: Creates a VPC and a public subnet, then exports their IDs.
```

Resources:

MyVPC:

Type: AWS::EC2::VPC

Properties:

CidrBlock: 10.0.0.0/16

EnableDnsSupport: true

EnableDnsHostnames: true

Tags:

- Key: Name

Value: MySharedVPC

MyPublicSubnet:

Type: AWS::EC2::Subnet

Properties:

VpcId: !Ref MyVPC

CidrBlock: 10.0.1.0/24

MapPublicIpOnLaunch: true

Tags:

- Key: Name

Value: MySharedPublicSubnet

MyInternetGateway:

Type: AWS::EC2::InternetGateway

Properties:

Tags:

- Key: Name

Value: MyInternetGateway

AttachGateway:

Type: AWS::EC2::VPCEGatewayAttachment

Properties:

VpcId: !Ref MyVPC

InternetGatewayId: !Ref MyInternetGateway

PublicRouteTable:

Type: AWS::EC2::RouteTable

Properties:

VpcId: !Ref MyVPC

Tags:

- Key: Name

Value: MyPublicRouteTable

PublicRoute:

Type: AWS::EC2::Route

DependsOn: AttachGateway # Ensure IGW is attached before creating route

Properties:

RouteTableId: !Ref PublicRouteTable

DestinationCidrBlock: 0.0.0.0/0

GatewayId: !Ref MyInternetGateway

SubnetRouteTableAssociation:

Type: AWS::EC2::SubnetRouteTableAssociation

Properties:

SubnetId: !Ref MyPublicSubnet

RouteTableId: !Ref PublicRouteTable

Outputs:

Export the VPC ID with a unique name

VPCId:

Description: The ID of the VPC

Value: !Ref MyVPC

Export:

Name: !Sub "\${AWS::StackName}-VPCId" # Ensures a unique export name per stack

Export the Public Subnet ID with a unique name

PublicSubnetId:

Description: The ID of the public subnet

Value: !Ref MyPublicSubnet

Export:

Name: !Sub "\${AWS::StackName}-PublicSubnetId"

Key parts in network-stack.yaml:

- **Outputs section:** This is where you define values that other stacks can reference.
- **Export: Name::** This property is crucial. It gives a unique name to the output value that can be imported by other stacks.
 - We use !Sub "\${AWS::StackName}-VPCId" to create a unique export name that includes the name of the CloudFormation stack itself. This helps prevent naming conflicts if you deploy multiple instances of this network stack.

2. Application Stack (application-stack.yaml) - The Importer

This stack deploys an EC2 instance and **imports** the PublicSubnetId and VpcId from the network-stack.

YAML

application-stack.yaml

AWSTemplateFormatVersion: '2010-09-09'

Description: Deploys an EC2 instance into an existing network using cross-stack references.

Parameters:

Parameter to get the name of the network stack from which to import values

NetworkStackName:

Type: String

Description: The name of the CloudFormation stack that exports the network details.

Resources:

WebServerSecurityGroup:

Type: AWS::EC2::SecurityGroup

Properties:

GroupName: WebAppSG

GroupDescription: Enable HTTP access to web server

Import the VPC ID from the network stack

VpcId: !ImportValue !Sub "\${NetworkStackName}-VPCId"

SecurityGroupIngress:

- IpProtocol: tcp

FromPort: 80

ToPort: 80

CidrIp: 0.0.0.0/0

Tags:

- Key: Name

Value: WebServerSecurityGroup

WebServerInstance:

Type: AWS::EC2::Instance

Properties:

ImageId: ami-0abcdef1234567890 # Replace with a valid AMI ID for your region (e.g., Amazon Linux 2 AMI)

InstanceType: t2.micro

Import the Public Subnet ID from the network stack

SubnetId: !ImportValue !Sub "\${NetworkStackName}-PublicSubnetId"

SecurityGroupIds:

- !Ref WebServerSecurityGroup

Tags:

- Key: Name

Value: MyAppWebServer

Outputs:

WebServerPublicIp:

Description: The public IP address of the web server

Value: !GetAtt WebServerInstance.PublicIp

Key parts in application-stack.yaml:

- **Parameters section:** The NetworkStackName parameter allows you to specify which network stack to import from. This makes the application stack reusable across different environments or network deployments.
- **!ImportValue intrinsic function:** This function is used to retrieve an exported output value from another stack.

- `!ImportValue !Sub "${NetworkStackName}-VPCId"`: This line dynamically constructs the export name (e.g., `MyNetworkStack-VPCId`) using the `NetworkStackName` parameter and then imports the corresponding VPC ID.
- Similarly for `SubnetId`.

How to deploy:

1. **Deploy network-stack.yaml first.** Give it a name, for example, `MyNetworkStack`. CloudFormation will create the VPC, subnet, etc., and then export the `VPCId` and `PublicSubnetId` with names like `MyNetworkStack-VPCId` and `MyNetworkStack-PublicSubnetId`.
2. **Deploy application-stack.yaml next.** When prompted for the `NetworkStackName` parameter, enter `MyNetworkStack` (or whatever name you used in step 1). The application stack will then import the VPC and subnet IDs from `MyNetworkStack` and use them to deploy the EC2 instance.

Advantages of Cross-Stack References:

- **Decoupling:** Stacks can be developed and managed independently. Changes to the network stack don't necessarily require changes to the application stack (unless the exported values themselves change).
- **Reusability:** The network stack can be a generic "base infrastructure" that multiple application stacks can use.
- **Modularity:** Breaks down large, complex infrastructure into smaller, more manageable units.
- **Clear Dependencies:** It's explicit which stacks depend on outputs from other stacks.

Important Considerations:

- **Region and Account:** Exported values can only be imported by stacks within the *same AWS account and AWS Region*.
- **Deletion/Updates:** You cannot delete or modify an exported output value in a stack if another stack is currently importing it. You must first update the importing stack(s) to remove the dependency before you can delete or change the exporting stack.
- **Uniqueness:** Export names must be unique within an AWS account and Region. Using `!Sub "${AWS::StackName}-..."` is a common best practice to ensure uniqueness.

1. What is a nested stack in AWS CloudFormation?

- A. A stack that is deployed across multiple regions
- B. A separate AWS account stack
- C. A stack created within another stack using the `AWS::CloudFormation::Stack` resource
- D. A stack that automatically rolls back upon failure

Answer: C

Why would you use a nested stack?

- A. To enable multi-region deployments
- B. To combine all resources into a single large template
- C. To modularize infrastructure and reuse templates
- D. To execute custom Lambda functions during stack creation

Answer: C

What is a cross-stack reference in CloudFormation?

- A. A reference to a stack deployed in another AWS region
- B. A way to export values from one stack and import into another
- C. A way to replicate stacks across accounts
- D. A method to store parameters in SSM

Answer: B

What keyword is used in CloudFormation to export a value from a stack?

- A. OutputShare
- B. Expose
- C. Export
- D. PublicRef

Answer: C

Which function is used in a dependent stack to bring in values from another stack?

- A. !Ref
- B. !GetAtt
- C. !IncludeValue
- D. !ImportValue

Answer: D