



## 1. What is AWS Fargate and what problems does it solve?

### Answer:

AWS Fargate is a **serverless compute engine** for containers that works with **Amazon ECS and Amazon EKS**. It eliminates the need to provision, configure, or manage EC2 instances.

### Problems it solves:

- **No server management:** You don't have to manage EC2 infrastructure.
- **Right-sizing:** You allocate only the required CPU/memory—no need to overprovision.
- **Scalability & security:** Each task/pod runs in its own kernel-isolated environment, improving security.

---

## 2. How does Fargate pricing work compared to EC2 instances? What kind of workload you will go for Fargate ?

### Answer:

Fargate charges **per second** based on the **CPU, memory, and runtime duration** of your tasks or pods.

### ◆ Fargate vs EC2:

Metric	Fargate	EC2
Pricing Model	Pay-per-use (CPU & memory secs)	Pay for full instance time
Management	Fully managed	Self-managed
Cost Efficiency	Great for short, bursty workloads	Better for long-running apps

It is like  
RENT vs BUY  
Decision

Use **EC2** if you have:

- Constant workloads
- Reserved instances
- Need for custom AMIs

Use **Fargate** if you want:

- Simplicity

- Auto-scaling without infra headaches

### 3. What are the supported CPU and memory configurations in Fargate?

Answer:

Fargate offers **specific combinations** of vCPU and memory. As of now:

- **CPU:** 0.25 vCPU to 16 vCPU
- **Memory:** 0.5 GB to 120 GB
- **Rules:**
  - Memory must be at least **2× vCPU** and at most **4× vCPU**



Example configurations:

#### vCPU Min Memory Max Memory

0.25	0.5 GB	2 GB
1	2 GB	4 GB
4	8 GB	16 GB

Use [AWS Fargate task size documentation](#) for full ranges.

### 4. Can you explain the difference between Fargate for ECS and Fargate for EKS?

Answer:

Feature	Fargate for ECS	Fargate for EKS
Orchestration	Amazon ECS	Amazon EKS (Kubernetes)
Abstraction	AWS-native abstraction	Kubernetes-native abstraction
Task type	ECS Tasks	Kubernetes Pods
Use case	Simpler, tightly-integrated apps	Complex, K8s-based apps

→ **Fargate for ECS** is simpler to set up and use.

→ **Fargate for EKS** offers more flexibility but requires knowledge of Kubernetes YAML specs and networking setup.

### 5. What are the limitations of AWS Fargate?

Answer:

While Fargate simplifies deployment, it comes with **limitations**:

- **No privileged containers** (e.g., Docker-in-Docker not allowed) (In GitHub Actions, GitLab CI, or AWS CodeBuild:  
If you try to run docker build inside a container job, it may **fail** or **hang**, because:  
The runner environment doesn't allow --privileged  
Docker daemon isn't accessible inside your container  
)
- **No EFS support in EKS Fargate** (only ECS supports EFS volumes)
- **Slower cold start** time vs EC2
- **Limited control over networking:** No custom ENIs per container
- **Can be more expensive** than EC2 for high-throughput apps
- **No GPU support** in Fargate
- **No daemon sets support in EKS Fargate**
- **Limited logging/monitoring customization**

## DEMO FARGET

### ✓ Overview

We'll deploy a simple Dockerized web app (e.g., Nginx) using Amazon ECS with Fargate launch type. This walkthrough covers:

- How Fargate abstracts away EC2
- Pricing per task (vs EC2)
- Supported CPU/memory
- Fargate for ECS vs EKS
- Fargate limitations in real-time

---

### 🔧 Prerequisites

- AWS account
- AWS CLI configured
- Docker installed (for building the image)
- IAM permissions for ECS, ECR, CloudWatch, IAM, VPC

---

## Deploy an NGINX Docker App on AWS Fargate

---

### Step 1: Create a Docker Image & Push to AWS ECR

#### 🔑 Goal:

Package your app into a Docker image and upload it to AWS Elastic Container Registry (ECR).

#### 👉 Steps:

1. **Create a folder for your app:**

```
bash
```

```
-----
```

```
mkdir fargate-demo && cd fargate-demo
```

2. **Add a basic Dockerfile using NGINX:**

```
bash
```

```
-----
```

```
echo "FROM nginx:alpine" > Dockerfile
```

3. **Build the image:**

```
bash
```

```
-----
```

```
docker build -t fargate-demo .
```

4. **Create an ECR repository in AWS to store this image:**

```
bash
```

```
-----
```

```
aws ecr create-repository --repository-name fargate-demo
```

✓ This will return output like:

```
json
```

```
-----
```

```
"repositoryUri": "906253564515.dkr.ecr.us-east-2.amazonaws.com/fargate-demo"
```

5. **Tag the Docker image to point to your ECR:**

```
bash
```

```
-----
```

```
aws_account_id=$(aws sts get-caller-identity --query Account --output text)
region=$(aws configure get region)
docker tag fargate-demo:latest
"$aws_account_id.dkr.ecr.$region.amazonaws.com/fargate-demo:latest"
```

#### 6. Login to ECR and push the image:

```
bash
```

```
-----
```

```
aws ecr get-login-password | docker login --username AWS --password-stdin
"$aws_account_id.dkr.ecr.$region.amazonaws.com"
docker push "$aws_account_id.dkr.ecr.$region.amazonaws.com/fargate-demo:latest"
```



### Step 2: Create ECS Cluster



#### Goal:

Create a place where your app will run inside ECS.



#### Command:

```
bash
```

```
-----
```

```
aws ecs create-cluster --cluster-name fargate-demo-cluster
```



Output will show your cluster details:

```
json
```

```
-----
```

```
"clusterName": "fargate-demo-cluster", "status": "ACTIVE"
```



### Step 3: Define Task Definition



#### Goal:

Tell ECS **how to run** your container (CPU, memory, image, port, etc.)



#### Create a file called task-definition.json with:

```
json
```

```
-----
```

```
{
  "family": "fargate-demo-task",
  "requiresCompatibilities": ["FARGATE"],
  "networkMode": "awsvpc",
  "cpu": "256",
  "memory": "512",
  "executionRoleArn": "arn:aws:iam::<ACCOUNT_ID>:role/ecsTaskExecutionRole",
  "containerDefinitions": [
    {
      "name": "fargate-demo-container",
      "image": "<ACCOUNT_ID>.dkr.ecr.<REGION>.amazonaws.com/fargate-demo:latest",
      "portMappings": [
        {
          "containerPort": 80,
          "protocol": "tcp"
        }
      ]
    }
  ]
}
```

```
    },
    "essential": true
  }
]
}
```

Replace <ACCOUNT\_ID> and <REGION> with your actual values.

---

#### Step 4: Register Task Definition

##### Goal:


Let ECS know about the task definition you created.

##### Command:

```
bash
```

```
-----
```

```
aws ecs register-task-definition --cli-input-json file://task-definition.json
```

 Output shows successful registration:

```
json
```

```
-----
```

```
"taskDefinitionArn": "arn:aws:ecs:us-east-2:906253564515:task-definition/fargate-demo-task:1"
```

---

#### Step 5: Create Security Group and Run the Task

##### Goal:

Allow web access and run the container.

##### Steps:

###### 1. Get default VPC and subnet:

```
bash
```

```
-----
```

```
vpc_id=$(aws ec2 describe-vpcs --filters Name=isDefault,Values=true --query "Vpcs[0].VpcId" --output text)
```

```
subnet_id=$(aws ec2 describe-subnets --filters Name=vpc-id,Values=$vpc_id --query "Subnets[0].SubnetId" --output text)
```

###### 2. Create security group and allow port 80 (HTTP):

```
bash
```

```
-----
```

```
sg_id=$(aws ec2 create-security-group --group-name fargate-sg --description "Fargate SG" --vpc-id $vpc_id --query "GroupId" --output text)
```

```
aws ec2 authorize-security-group-ingress --group-id $sg_id --protocol tcp --port 80 --cidr 0.0.0.0/0
```

###### 3. Run the Fargate task:

```
bash
```

```
-----
```

```
aws ecs run-task \
  --cluster fargate-demo-cluster \
  --launch-type FARGATE \
```

```
--network-configuration
"awsvpcConfiguration={subnets=[${subnet_id}],securityGroups=[${sg_id}],assignPublicIp=
ENABLED}" \
--task-definition fargate-demo-task
```

---

## Step 6: Access Your App in Browser

### Goal:

Get the **public IP** of the running container and open it in a browser.

### Steps:

#### 1. Get the running task ARN:

```
bash
```

```
-----
```

```
task_arn=$(aws ecs list-tasks --cluster fargate-demo-cluster --query "taskArns[0]" --
output text)
```

#### 2. Get the network interface (ENI):

```
bash
```

```
eni_id=$(aws ecs describe-tasks --cluster fargate-demo-cluster --tasks $task_arn --
query "tasks[0].attachments[0].details[?name=='networkInterfaceId'].value" --output
text)
```

#### 3. Get the public IP:

```
bash
```

```
-----
```

```
aws ec2 describe-network-interfaces --network-interface-ids $eni_id --query
"NetworkInterfaces[0].Association.PublicIp" --output text
```

#### 4. Open in browser:

```
cpp
```

```
-----
```

```
http://<public-ip>
```

You'll see the NGINX welcome page!

---

 Not Secure 18.119.116.6

## Welcome to nginx!


If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to [nginx.org](https://nginx.org).  
Commercial support is available at [nginx.com](https://nginx.com).

*Thank you for using nginx.*


**What is AWS Fargate primarily used for?**

- A. Hosting serverless web apps
- B. Running containers without managing servers
- C. Running Lambda functions
- D. Managing EC2 Auto Scaling groups

B. Running containers without managing servers 

**Which of the following is NOT a benefit of AWS Fargate?**

- A. Eliminates need for EC2 instance management
- B. Enables overprovisioning for performance
- C. Improves security through kernel isolation
- D. Allocates just the required CPU and memory

C. Improves security through kernel isolation 

**For which type of workload is AWS Fargate most cost-efficient?**

- A. Long-running batch jobs
- B. Constant 24x7 workloads
- C. Short, bursty, or unpredictable workloads
- D. High-performance computing with custom kernels

C. Short, bursty, or unpredictable workloads 

**Which of the following is a limitation of AWS Fargate?**

- A. Can't scale automatically
- B. Doesn't support Kubernetes
- C. Limited CPU/memory configuration options
- D. Lacks integration with ECS

C. Limited CPU/memory configuration options 