



6. How do you implement persistent storage with Fargate?

Answer:

Fargate tasks are ephemeral by default, which means data is lost once the task stops. To persist data:

- I use **EFS (Elastic File System)** – a managed, scalable file system by AWS.
- Steps:
 1. Create an **EFS file system** in the same VPC.
 2. Mount it to the Fargate task using **task definition volumes**.
 3. Define the mount point inside the container so that any data written there persists.
- This is useful for shared or durable storage like logs, user uploads, or session data.

◆ Step 1: Create an EFS File System

1. Go to the **Amazon EFS Console**.
2. Click **Create file system**.
3. Choose the **same VPC** as your Fargate service.
4. Under **Access points**, optionally create one (recommended for managed permissions).
5. Add mount targets in the **same subnets and AZs** where your Fargate tasks run.
6. Ensure your **EFS security group** allows inbound NFS traffic (port 2049) from Fargate security group.

◆ Step 2: Task Execution Role (IAM Permissions)

Ensure the task execution role has the following policy:

json

```
{
  "Effect": "Allow",
  "Action": [
    "elasticfilesystem:ClientMount",
    "elasticfilesystem:ClientWrite",
    "elasticfilesystem:DescribeMountTargets"
  ],
  "Resource": "*"
}
```

```
}
```

◆ Step 3: Create ECS Task Definition with EFS Volume

Here is a sample **Fargate ECS task definition (JSON)** with EFS:

json

```
{
  "family": "fargate-efs-demo",
  "networkMode": "awsvpc",
  "requiresCompatibilities": ["FARGATE"],
  "cpu": "256",
  "memory": "512",
  "executionRoleArn": "arn:aws:iam::<account-id>:role/ecsTaskExecutionRole",
  "containerDefinitions": [
    {
      "name": "app",
      "image": "nginx",
      "portMappings": [
        {
          "containerPort": 80,
          "protocol": "tcp"
        }
      ],
      "mountPoints": [
        {
          "sourceVolume": "efs-volume",
          "containerPath": "/usr/share/nginx/html", // your mount path
          "readOnly": false
        }
      ],
      "logConfiguration": {
        "logDriver": "awslogs",
        "options": {
          "awslogs-group": "/ecs/efs-demo",
          "awslogs-region": "us-east-1",
          "awslogs-stream-prefix": "ecs"
        }
      }
    }
  ],
  "volumes": [
    {
      "name": "efs-volume",
      "efsVolumeConfiguration": {
        "fileSystemId": "fs-12345678", // replace with your EFS ID
        "transitEncryption": "ENABLED",
        "authorizationConfig": {
```

```
"accessPointId": "fsap-abcdefgh123456789", // optional but recommended
"iam": "ENABLED"
}
}
}
]
```

◆ Step 4: Launch Fargate Task

- Make sure you use:
 - Same **VPC** and **subnet** as EFS.
 - Correct **security groups** that allow NFS (port 2049) from ECS task ENIs.

◆ Step 5: Verify

- Access the container and create/write files to the mounted directory.
- Stop and re-run the task. Files will persist.

7. Explain how networking works in Fargate tasks.

Answer:

Fargate uses **awsvpc mode** for networking. That means:

- Each task gets its own **ENI (Elastic Network Interface)** and **private IP** address.
- It behaves like a separate EC2 instance within the **VPC**.
- You can control traffic using:
 - **Subnets** (public/private),
 - **Security Groups** (like firewalls),
 - **Route tables** and **NACLs**.
- For internet access:
 - Public subnet + public IP, or
 - Private subnet + NAT Gateway.

This makes Fargate secure, as each task is isolated at the network level.

8. How do you troubleshoot Fargate tasks that fail to start?

Answer:

I follow this checklist to debug Fargate tasks:

1. **Check CloudWatch Logs:**
 - Look for stdout or stderr logs from the container.
2. **Describe Task:**
 - Run `aws ecs describe-tasks` to see error messages (e.g., `RESOURCE:MEMORY`, `CannotPullContainerError`).
3. **IAM Role Issues:**
 - Ensure the task role and execution role have proper permissions (e.g., to pull from ECR or write to CloudWatch).
4. **Image Pull Failures:**
 - Verify the image exists in ECR or Docker Hub and credentials are correct.
5. **Check Networking:**

- Make sure the security groups and subnets allow necessary communication (e.g., DB ports).

9. What are the best practices for optimizing Fargate costs?

Answer:

To reduce Fargate costs:

- **Right-size tasks:** Avoid over-provisioning CPU/memory.
- **Use Spot Fargate** for non-critical workloads (~70% cheaper).
- **Task Auto Scaling:** Use ECS Service Auto Scaling to run only what you need.
- **Use Graviton-based compute:** If supported, they're cheaper and more efficient.
- **Shut down idle tasks:** Schedule cleanup of unused resources.
- **Monitor usage:** Set CloudWatch alarms for unexpected cost spikes.

10. How do you implement logging and monitoring for Fargate tasks?

Answer:

Logging:

- Enable **CloudWatch logging** in task definition:
 - Add a **logConfiguration** block to use the awslogs driver.
 - Each task sends logs to a specific log group.

Monitoring:

- Use **CloudWatch metrics** to track:
 - CPU, memory usage, task health.
- Use **CloudWatch Alarms** to alert on high usage or failures.
- For more insights:
 - Integrate with **AWS X-Ray** or **Datadog/Grafana/Prometheus** if needed.

11. Explain how to use Fargate with VPC endpoints for private connectivity.

Answer:

If I want my Fargate task to securely access AWS services **without using the public internet**, I use **VPC Endpoints**:

- **Step-by-step:**
 1. Create a **VPC Endpoint** (e.g., for S3, DynamoDB, or ECR).
 2. Make sure Fargate tasks run in **private subnets**.
 3. Attach proper **IAM roles** and **route tables**.
 4. The task now communicates directly over AWS's internal network – **no need for NAT Gateway or internet gateway**.

This is important for **secure, low-latency access** to AWS services in private environments.

Deploy an NGINX Docker App on AWS Fargate

Step 1: Create a Docker Image & Push to AWS ECR



Goal:

Package your app into a Docker image and upload it to AWS Elastic Container Registry (ECR).



Steps:

1. **Create a folder for your app:**

```
bash
```

```
-----
```

```
mkdir fargate-demo && cd fargate-demo
```

2. Add a basic Dockerfile using NGINX:

```
bash
```

```
-----
```

```
echo "FROM nginx:alpine" > Dockerfile
```

3. Build the image:

```
bash
```

```
-----
```

```
docker build -t fargate-demo .
```

4. Create an ECR repository in AWS to store this image:

```
bash
```

```
-----
```

```
aws ecr create-repository --repository-name fargate-demo
```

✅ This will return output like:

```
json
```

```
-----
```

```
"repositoryUri": "906253564515.dkr.ecr.us-east-2.amazonaws.com/fargate-demo"
```

5. Tag the Docker image to point to your ECR:

```
bash
```

```
-----
```

```
aws_account_id=$(aws sts get-caller-identity --query Account --output text)
```

```
region=$(aws configure get region)
```

```
docker tag fargate-demo:latest
```

```
"$aws_account_id.dkr.ecr.$region.amazonaws.com/fargate-demo:latest"
```

6. Login to ECR and push the image:

```
bash
```

```
-----
```

```
aws ecr get-login-password | docker login --username AWS --password-stdin
```

```
"$aws_account_id.dkr.ecr.$region.amazonaws.com"
```

```
docker push "$aws_account_id.dkr.ecr.$region.amazonaws.com/fargate-demo:latest"
```



Step 2: Create ECS Cluster



Goal:

Create a place where your app will run inside ECS.



Command:

```
bash
```

```
-----
```

```
aws ecs create-cluster --cluster-name fargate-demo-cluster
```

✅ Output will show your cluster details:

```
json
```

```
-----
```

```
"clusterName": "fargate-demo-cluster", "status": "ACTIVE"
```



Step 3: Define Task Definition



Goal:

Tell ECS **how to run** your container (CPU, memory, image, port, etc.)

 **Create a file called task-definition.json with:**

json

```
{
  "family": "fargate-demo-task",
  "requiresCompatibilities": ["FARGATE"],
  "networkMode": "awsvpc",
  "cpu": "256",
  "memory": "512",
  "executionRoleArn": "arn:aws:iam::<ACCOUNT_ID>:role/ecsTaskExecutionRole",
  "containerDefinitions": [
    {
      "name": "fargate-demo-container",
      "image": "<ACCOUNT_ID>.dkr.ecr.<REGION>.amazonaws.com/fargate-demo:latest",
      "portMappings": [
        {
          "containerPort": 80,
          "protocol": "tcp"
        }
      ],
      "essential": true
    }
  ]
}
```

Replace <ACCOUNT_ID> and <REGION> with your actual values.

Step 4: Register Task Definition

Goal:

Let ECS know about the task definition you created.

 **Command:**

bash

```
aws ecs register-task-definition --cli-input-json file://task-definition.json
```

✅ Output shows successful registration:

json

```
"taskDefinitionArn": "arn:aws:ecs:us-east-2:906253564515:task-definition/fargate-
demo-task:1"
```

Step 5: Create Security Group and Run the Task

Goal:

Allow web access and run the container.

 **Steps:**

1. **Get default VPC and subnet:**

bash

```
vpc_id=$(aws ec2 describe-vpcs --filters Name=isDefault,Values=true --query
"Vpcs[0].VpcId" --output text)
subnet_id=$(aws ec2 describe-subnets --filters Name=vpc-id,Values=$vpc_id --query
"Subnets[0].SubnetId" --output text)
```

2. Create security group and allow port 80 (HTTP):

```
bash
```

```
-----
```

```
sg_id=$(aws ec2 create-security-group --group-name fargate-sg --description "Fargate
SG" --vpc-id $vpc_id --query "GroupId" --output text)
```

```
aws ec2 authorize-security-group-ingress --group-id $sg_id --protocol tcp --port 80 --
cidr 0.0.0.0/0
```

3. Run the Fargate task:

```
bash
```

```
-----
```

```
aws ecs run-task \
  --cluster fargate-demo-cluster \
  --launch-type FARGATE \
  --network-configuration
"awsvpcConfiguration={subnets=[$subnet_id],securityGroups=[$sg_id],assignPublicIp=
ENABLED}" \
  --task-definition fargate-demo-task
```

Step 6: Access Your App in Browser



Goal:

Get the **public IP** of the running container and open it in a browser.



Steps:

1. Get the running task ARN:

```
bash
```

```
-----
```

```
task_arn=$(aws ecs list-tasks --cluster fargate-demo-cluster --query "taskArns[0]" --
output text)
```

2. Get the network interface (ENI):

```
bash
```

```
eni_id=$(aws ecs describe-tasks --cluster fargate-demo-cluster --tasks $task_arn --
query "tasks[0].attachments[0].details[?name=='networkInterfaceId'].value" --output
text)
```

3. Get the public IP:

```
bash
```

```
-----
```

```
aws ec2 describe-network-interfaces --network-interface-ids $eni_id --query
"NetworkInterfaces[0].Association.PublicIp" --output text
```

4. Open in browser:

```
cpp
```

```
-----
```

http://<public-ip>

You'll see the NGINX welcome page!

⚠ Not Secure 18.119.116.6

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

How can you implement persistent storage for a Fargate task?

- A. Attach EBS volumes directly to the Fargate task
- B. Use S3 to mount as a filesystem
- C. Use EFS and mount it via task definition
- D. Enable data persistence in task settings

✅ **Correct Answer:** C. Use EFS and mount it via task definition

Q. Which networking mode is used by AWS Fargate for tasks?

- A. bridge
- B. host
- C. awsvpc
- D. overlay

✅ **Correct Answer:** C. awsvpc