@devopschallengehub







EKS Deployment Flow

-

STEP 1: Create Flask App

Files:

- app.py
- requirements.txt

Code:

from flask import Flask
app = Flask(__name__)
@app.route("/")
def hello():
 return "Hello World from EKS"

Command:

> python app.py

Purpose: Have a working app before deployment.

-

STEP 2: Containerize the App

File: Dockerfile

Dockerfile:

FROM python:3.9-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install -r requirements.txt
COPY app.py .
CMD ["python", "app.py"]

Commands:

> docker build -t flask-hello-world .

> docker run -p 5000:5000 flask-hello-world

Purpose: Package app into a container.

_

STEP 3: Push to Amazon ECR

Commands:

- > aws ecr create-repository --repository-name flask-hello-world
- > aws ecr get-login-password | docker login ...
- > docker tag flask-hello-world <ecr-uri>
- > docker push <ecr-uri>

Purpose: Store image on AWS for use in EKS. yaml

_

STEP 4: Create EKS Cluster

File: cluster-config.yaml

YAML:

kind: ClusterConfig

metadata:

name: flask-demo-cluster

region: us-west-2

nodeGroups:

- name: worker-nodes instanceType: t3.medium

Command:

> eksctl create cluster -f cluster-config.yaml

Purpose: Provision Kubernetes control plane + nodes.

-

_

STEP 5: Add Storage with PVC

Files: storage-class.yaml, pvc.yaml

storage-class.yaml: kind: StorageClass

provisioner: ebs.csi.aws.com

pvc.yaml:

kind: PersistentVolumeClaim

spec:

resources: requests: storage: 1Gi

Commands:

- > kubectl apply -f storage-class.yaml
- > kubectl apply -f pvc.yaml

Purpose: Provide durable disk space to your app.

yaml

🖋 STEP 6: Deploy Flask App

File: flask-deployment.yaml

YAML (simplified):

kind: Deployment

spec:

replicas: 2 containers:

- name: flask-app

image: <your-ecr-image>

volumeMounts:

 mountPath: /app/data name: flask-storage

volumes:

 name: flask-storage persistentVolumeClaim: claimName: flask-storage

Command:

> kubectl apply -f flask-deployment.yaml

Purpose: Start app pods in EKS using container image.

yaml

STEP 7: Create Internal Service

File: flask-service.yaml

YAML:

kind: Service

spec:

selector:

app: flask-app

ports:

- port: 80

targetPort: 5000 type: ClusterIP

Command:

> kubectl apply -f flask-service.yaml

Purpose: Access pods inside the cluster. yaml

_

STEP 8: Expose App to Internet

File: flask-ingress.yaml

YAML:

kind: Ingress metadata: annotations:

alb.ingress.kubernetes.io/scheme: internet-facing

spec:

rules: - http:

paths:

- path: / backend: service:

name: flask-service

port:

number: 80

Command:

> kubectl apply -f flask-ingress.yaml

> kubectl get ingress

Purpose: Give public access via AWS ALB. yaml

,

STEP 9: Secure Access (RBAC + IAM)

File: rbac.yaml

YAML: kind: Role rules:

- apiGroups: [""]
resources: ["pods"]
verbs: ["get", "list"]

kind: RoleBinding

subjects:

- kind: ServiceAccount

name: flask-sa

roleRef: kind: Role

name: flask-role

Commands:

> kubectl apply -f rbac.yaml

Purpose: Control what your app or users can access.

yaml

STEP 10: Harden Security in Deployment

File: flask-deployment-secure.yaml

YAML:

securityContext: runAsNonRoot: true runAsUser: 1000 containers:

- securityContext:

allowPrivilegeEscalation: false readOnlyRootFilesystem: true

Purpose: Prevent privilege abuse, follow security best practices.

markdown

STEP 11: Verify Everything

Commands:

- > kubectl get nodes
- > kubectl get pods
- > kubectl logs <pod-name>
- > kubectl get svc
- > kubectl get ingress
- > curl <ALB-URL>

Purpose: Ensure everything is running and accessible.

Summary

Step Purpose Key File / Command 1 Write app app.py 2 Containerize Dockerfile, docker build 3 Push to ECR docker push, aws ecr 4 Create EKS Cluster cluster-config.yaml, eksctl create Add Storage storage-class.yaml, pvc.yaml 5 6 Deploy app flask-deployment.yaml

Step Purpose Key File / Command

Expose internally flask-service.yaml
 Expose externally (public) flask-ingress.yaml

9 Add access control rbac.yaml, kubectl apply

10 Harden container runtime securityContext in deployment YAML

11 Test and verify kubectl, curl, browser

1. How do you implement RBAC (Role-Based Access Control) in EKS?

RBAC = "Who can do what?"

RBAC stands for **Role-Based Access Control**, a system used in **Kubernetes** (and therefore in EKS too) to control **who can access what** resources in the cluster. For example:

- Dev users can only view Pods.
- Admins can create or delete resources.
- CI/CD pipelines may get special access to certain namespaces.

EKS + RBAC = IAM + Kubernetes RBAC

In Amazon EKS, you manage access at two levels:

- IAM level: Controls which users/roles from AWS can authenticate into the cluster.
- 2. **Kubernetes RBAC level**: Controls **what those users can do** after logging in. So even if an IAM user is allowed to log in to the cluster, RBAC decides what they can actually do inside Kubernetes.

How IAM and RBAC Work Together in EKS

In EKS:

Layer Controls Example

IAM (AWS) Who can authenticate to the cluster IAM user kube_test_user

RBAC (K8s) What that user can do inside Kubernetes Role allowing list/get pods

IAM = Authentication

RBAC = Authorization

Example: Grant kube_test_user read-only access to pods

□AM User Setup (Assumed Already Exists)

You already have this IAM user:

ruby

Name: kube_test_user

ARN: arn:aws:iam::<account-id>:user/kube_test_user

Make sure the IAM user has permission to call eks:DescribeCluster and sts:AssumeRole.

21 Update aws-auth ConfigMap to Allow IAM Access

You must tell EKS:

"This IAM user maps to a Kubernetes user".

Run: bash

kubectl edit configmap aws-auth -n kube-system

Add under mapUsers::

yaml

mapUsers:

- userarn: arn:aws:iam::<account-id>:user/kube_test_user

username: kube_test_user

groups:

- read-only-group

This says: "Let kube_test_user authenticate, and map them to the Kubernetes group

read-only-group."

₹Create an RBAC Role for the Group

Create file: rbac-readonly.yaml

yaml

apiVersion: rbac.authorization.k8s.io/v1

kind: Role metadata:

name: readonly-pods namespace: demo

rules:

- apiGroups: [""]
resources: ["pods"]

verbs: ["get", "list", "watch"]

apiVersion: rbac.authorization.k8s.io/v1

kind: RoleBinding

metadata:

name: bind-readonly-pods

namespace: demo

subjects:

- kind: Group

name: read-only-group # This matches the group in aws-auth

apiGroup: rbac.authorization.k8s.io

roleRef: kind: Role

name: readonly-pods

apiGroup: rbac.authorization.k8s.io

Apply it: bash

kubectl apply -f rbac-readonly.yaml

⚠Test Access (From kube_test_user CLI)

From a system where the kube_test_user is configured in AWS CLI:

bash

aws eks update-kubeconfig --name your-cluster-name --region your-region \ --role-arn arn:aws:iam::<account-id>:user/kube_test_user

Try this:

bash

kubectl get pods -n demo



This should work.

Try this:

bash

kubectl delete pod mypod -n demo



This should fail – because the role is read-only.

Real-World Use Case

- You want developers to view Pods but not delete anything.
- You want **DevOps engineers** to manage deployments only in dev namespace.
- You want CI/CD systems to have scoped permissions, not full admin rights.

2. What is AWS IAM integration with EKS?

Amazon EKS (Elastic Kubernetes Service) runs Kubernetes clusters on AWS.

To secure it, we need to:

- 1. Control who can access the cluster (admins, devs, CI/CD users).
- 2. Control what running Pods can access in AWS (like S3, DynamoDB, etc.).

To do this, EKS integrates with AWS IAM (Identity and Access Management).

Two Key Parts of IAM Integration in EKS

Purpose Technique Use Case Cluster access Map IAM users/roles to Kubernetes aws-auth ConfigMap (users) users Pod access to AWS IAM Roles for Service Let apps in pods access S3, services Accounts (IRSA) DynamoDB, etc. securely

Part 1: IAM → Kubernetes Access (via aws-auth ConfigMap) What it does: Maps IAM Users or Roles to Kubernetes RBAC identities. Where? In a special ConfigMap inside the EKS cluster: bash kubectl edit configmap aws-auth -n kube-system **Example:** vaml mapUsers: - userarn: arn:aws:iam::123456789012:user/dev username: dev-user groups: - system:masters This allows: IAM user dev to connect to the EKS cluster as dev-user Grants them full access (system:masters) unless you restrict with RBAC Check access: bash kubectl auth can-i list pods --as dev-user Part 2: IAM Roles for Service Accounts (IRSA) 🖑 Problem: Let's say a pod needs to access an S3 bucket. How do we give it access safely, without hardcoding AWS credentials? Solution: Use IRSA (IAM Roles for Service Accounts) IRSA = Secure way to attach an IAM Role to a Kubernetes ServiceAccount. So, the Pod assumes that IAM Role, and gets permissions to AWS services — no hardcoded creds. How IRSA Works 1. You create an IAM Role with the needed AWS permissions. 2. Attach an OIDC trust policy, so EKS knows this role can be assumed by a 3. In Kubernetes, you create a **ServiceAccount** linked to that IAM Role. 4. Your **Pod uses that ServiceAccount**, and it gets AWS access automatically.

Example YAML for IRSA: IAM Role trust policy: json ------

"Effect": "Allow",

```
"Principal": {
 "Federated": "arn:aws:iam::123456789012:oidc-provider/oidc.eks.us-west-
2.amazonaws.com/id/EXAMPLED539D4633E53DE1B716D3041E"
"Action": "sts:AssumeRoleWithWebIdentity",
"Condition": {
 "StringEquals": {
  "oidc.eks...:sub": "system:serviceaccount:default:myapp-sa"
 }
}
Kubernetes ServiceAccount:
yaml
apiVersion: v1
kind: ServiceAccount
metadata:
name: myapp-sa
annotations:
 eks.amazonaws.com/role-arn: arn:aws:iam::123456789012:role/MyAppIRSA
Then in your Deployment/Pod spec:
yaml
serviceAccountName: myapp-sa
Now the pod can access AWS services like S3 or DynamoDB with fine-grained
permissions.
```

Summary for Beginners

Feature Purpose Where It Works

aws-auth Maps IAM users/roles to Controls who can access the EKS

ConfigMap K8s users cluster

IRSA Maps IAM roles to Pods Controls what apps inside EKS can access in AWS



- A developer logs in using their IAM user to view resources: aws-auth mapping.
- A Pod uploads files to S3: IRSA.
- A **CI/CD pipeline** (e.g., GitHub Actions) assumes an IAM role and deploys to EKS.

.

3. How do you configure service accounts in EKS?

a First: What's the Problem We're Solving?

Imagine you have a **Pod in EKS** that needs to access:

- An S3 bucket to upload logs
- Or a **DynamoDB table** to read/write data

But — how does that Pod get AWS credentials securely?

We don't want to:

- Hardcode AWS keys in the Pod
- Use overly broad permissions

Solution: Use IAM Roles for Service Accounts (IRSA)

With IRSA, you attach a specific **IAM Role** to a **Kubernetes ServiceAccount**. Pods using that ServiceAccount will **automatically get AWS credentials** — securely and scoped.

Steps to Configure IRSA in EKS (Beginner Friendly)

Step 1: Enable IAM OIDC Provider

Each EKS cluster needs an **OIDC provider** (OpenID Connect) so IAM can trust Kubernetes.

Run this command:

bash

eksctl utils associate-iam-oidc-provider \

- --region <your-region> \
- --cluster <your-cluster-name> \
- --approve

This allows your IAM roles to trust tokens from your EKS cluster.

◆ Step 2: Create IAM Policy

Create a policy with least-privilege AWS permissions.

☑ Example: S3 Read/Write

Save it as s3-access-policy.json, then run:

```
bash
-----
aws iam create-policy \
--policy-name S3AccessPolicy \
--policy-document file://s3-access-policy.json
```

Step 3: Create IAM Role with Trust Policy

This IAM Role:

- Can only be assumed by EKS Pods
- Only when using a specific ServiceAccount

```
Use this trust policy (replace placeholders):
```

```
ison
-----
 "Version": "2012-10-17",
"Statement": [{
 "Effect": "Allow",
 "Principal": {
  "Federated": "arn:aws:iam::<account-id>:oidc-
provider/oidc.eks.<region>.amazonaws.com/id/<eks-id>"
 "Action": "sts:AssumeRoleWithWebIdentity",
 "Condition": {
  "StringEquals": {
   "oidc.eks.<region>.amazonaws.com/id/<eks-id>:sub":
"system:serviceaccount:default:s3-access-sa"
 }
}]
Then create the IAM role:
bash
aws iam create-role \
--role-name S3AccessRole \
--assume-role-policy-document file://trust-policy.json
And attach the S3 policy to it:
bash
aws iam attach-role-policy \
--role-name S3AccessRole \
--policy-arn arn:aws:iam::<account-id>:policy/S3AccessPolicy
```

Step 4: Create a Kubernetes ServiceAccount (with annotation)

This is the link between your pod and the IAM role.

Example YAML:

yaml

```
apiVersion: v1
kind: ServiceAccount
metadata:
name: s3-access-sa
annotations:
 eks.amazonaws.com/role-arn: arn:aws:iam::<account-id>:role/S3AccessRole
Apply it:
bash
kubectl apply -f sa.yaml
Step 5: Use the ServiceAccount in Your Pod/Deployment
Now update your Pod or Deployment YAML to use the service account.
Example:
yaml
apiVersion: apps/v1
kind: Deployment
metadata:
name: myapp
spec:
replicas: 1
selector:
 matchLabels:
  app: myapp
template:
 metadata:
  labels:
   app: myapp
 spec:
  serviceAccountName: s3-access-sa
  containers:
  - name: mycontainer
   image: myapp:latest
   # App can now access AWS S3 via SDK without hardcoding credentials
Why Is This Great?

    Secure: No need to store access keys in pods

   • © Fine-grained access: Each pod gets only the permissions it needs
   • Cloud-native: Works seamlessly with AWS SDKs and CLI in your app
Summary
Step What You Do
                                     Why
     Enable OIDC
                                     So IAM can trust EKS tokens
1□
```

Define what AWS resources are allowed

2

Create IAM Policy

Step What You Do		Why
3	Create IAM Role with Trust	Link the IAM Role to EKS ServiceAccount
40	Create Kubernetes ServiceAccount	Annotate it with IAM Role ARN
5	Use ServiceAccount in Pod	Your app gets secure AWS access automatically

4. What are Pod Security Standards and how do you implement them?

In simple terms:

PSS are **rules that control how secure your Pods are** — like what they're allowed to do, and what they're **not** allowed to do.

Kubernetes defines 3 security levels:

Level Purpose

privileged No restrictions. Good for debugging, but unsafe for production.

baseline Prevents known bad practices (e.g., running as root). Safer for most workloads.

 $\begin{tabular}{ll} \textbf{restricted} & \textbf{Very strict. Follows best practices for least privilege and hardening. Ideal for production.} \end{tabular}$

! Why Do You Need PSS in EKS?

Without any restrictions:

- Pods could run as root
- Mount host paths (accessing the EC2 machine itself!)
- Use **privileged mode** (full admin rights inside container)

This opens up security risks like container breakout attacks.

So, PSS lets you enforce rules at the namespace level to protect the cluster.

How to Implement Pod Security Standards (Step-by-Step)

Kubernetes 1.25+ (used in modern EKS versions) comes with a built-in **PodSecurity** admission controller.

Step 1: Choose Your Namespace

Let's say your namespace is dev.

You want all pods in dev to follow restricted security policies.

Step 2: Add Labels to Enforce PSS

Use this command:

bash

kubectl label namespace dev \
pod-security.kubernetes.io/enforce=restricted

This means: **Pods in the dev namespace must follow the restricted policy.**

You can also add optional labels for audit or warn:

bash

pod-security.kubernetes.io/warn=baseline pod-security.kubernetes.io/audit=privileged

So Kubernetes will warn or log if a pod violates less strict policies.

Example Label Setup:

Label Meaning

pod-security.kubernetes.io/enforce=restricted Block anything below restricted pod-security.kubernetes.io/warn=baseline Show warning if policy is below baseline pod-security.kubernetes.io/audit=privileged Log policy violations of privileged level

Example: What Gets Blocked?

If someone tries to apply this Pod YAML:

yaml

apiVersion: v1 kind: Pod metadata: name: badpod

spec:

containers:
- name: alpine
image: alpine
securityContext:
privileged: true

It will get **denied** in a namespace labeled restricted because:

Running a privileged container is not allowed under restricted mode.

Sonus: Advanced Policy with OPA Gatekeeper or Kyverno

If you want **custom rules** (e.g., "pods must use a specific image registry" or "no containers with hostPath"), use:

Tool Use Case

OPA Gatekeeper Policy as code using Rego language

Kyverno Simpler YAML-based policy engine, very Kubernetes-native

These tools provide **richer validation**, auditing, and auto-remediation.

Summary for Beginners

Concept Meaning

Pod Security Standards (PSS) Built-in security levels for Pods

privileged No restrictions (not safe)

baseline Medium safety — blocks known bad configs

Concept Meaning

restricted Strong safety — best for production

Namespace labels Used to enforce security levels

OPA/Gatekeeper/Kyverno Used for custom advanced policy enforcement

brace Real-World Use Case in EKS

- You run a multi-tenant EKS cluster (e.g., dev, staging, prod).
- You want to make sure developers can't run insecure containers.
- You label dev and prod namespaces with different PSS levels.
- In prod, only hardened, secure workloads can run.

5. How do you handle secrets management in EKS?

s What Are Secrets and Why Manage Them?

In applications, secrets = sensitive info like:

- Database passwords
- API tokens
- SSH keys
- AWS credentials

If these secrets are not stored securely, **anyone who compromises a pod** can steal them — leading to data leaks or system compromise.

So in EKS (Kubernetes on AWS), managing secrets safely is **critical**.

3 Options for Secrets Management in EKS

Let's go through them **from basic to advanced**, so beginners can understand and build confidence.

Option 1: Kubernetes Secrets (Basic)

- Secrets are stored inside Kubernetes using built-in Secret resources.
- They are base64-encoded, not encrypted by default.
- Stored in etcd, which should be encrypted (in EKS, it is by default).

Example YAML:

yaml

apiVersion: v1 kind: Secret metadata:

name: db-creds type: Opaque

data:

username: YWRtaW4= # "admin"

password: MWYyZDFlMmU2N2Rm # "1f2d1e2e67df"

To create it:

bash

kubectl apply -f secret.yaml

Then in a Pod:
yaml
----env:
- name: DB_USER
valueFrom:
secretKeyRef:
name: db-creds
key: username

Limitations:

- Not encrypted at rest (unless etcd encryption is configured).
- Secrets visible to anyone with access to the namespace.
- Better for non-critical dev/test use.

Option 2: AWS Secrets Manager with IRSA (Production-Ready)

AWS Secrets Manager is a fully managed secrets service:

- Secrets are encrypted with KMS.
- You can rotate secrets automatically.
- Access is controlled via IAM (and in EKS, via IRSA).

Flow:

- 1. Store your secret in AWS Secrets Manager.
- 2. Create a Kubernetes ServiceAccount with IRSA.
- 3. Pod uses AWS SDK (or init/sidecar container) to fetch the secret securely.

Example:

Store a secret:

bash

aws secretsmanager create-secret \

- --name db-creds \
- --secret-string '{"username":"admin","password":"s3cr3t"}'

In your app (Node.js, Python, etc.), use the AWS SDK:

is

const client = new SecretsManager();

const secret = await client.getSecretValue({ SecretId: 'db-creds' }).promise();

Best for production. Supports audit logging, rotation, and IAM-based fine control.

Option 3: Secrets Store CSI Driver (Recommended Integration)

This lets Kubernetes **mount external secrets** (like from Secrets Manager or Vault) as **volumes**, so your app doesn't need to call AWS SDK directly.

This is great if:

- You don't want to modify your app code.
- You want secret values injected directly as files.

How It Works:

- 1. Install **Secrets Store CSI Driver** in your EKS cluster.
- 2. Install AWS provider plugin.
- 3. Create a SecretProviderClass that tells which secrets to pull.

4. Mount those secrets in Pod using a CSI volume.

Example:

yaml

apiVersion: secrets-store.csi.x-k8s.io/v1

kind: SecretProviderClass

metadata:

name: aws-secrets

spec:

provider: aws parameters: objects: |

> - objectName: "db-creds" objectType: "secretsmanager"

Then in your pod:

yaml

volumeMounts:

- name: secrets-store

mountPath: "/mnt/secrets"

volumes:

- name: secrets-store

csi:

driver: secrets-store.csi.k8s.io

readOnly: true volumeAttributes:

secretProviderClass: "aws-secrets"

Now /mnt/secrets/db-creds contains the secret securely — no need for base64 decoding or SDK calls.

Summary for Beginners

Method	Description	Best For
Kubernetes Secrets	Built-in, simple, stored in etcd	Dev/test environments
AWS Secrets Manager + IRSA	Securely access AWS secrets via IAM roles	Production apps needing AWS secrets
Secrets Store CSI Driver	Mount secrets as volumes from AWS Secrets Manager/Vault	No-code change, production-grade

Recommendation for Beginners in EKS

- Start with Kubernetes Secrets in test/dev.
- Learn IRSA to fetch from AWS Secrets Manager in production.
- Use Secrets Store CSI Driver for secure volume-based access especially useful in regulated environments.

6. What are the security best practices for EKS?

Overview: Why Security in EKS Matters

EKS (Elastic Kubernetes Service) gives you the power of Kubernetes on AWS.

But with great power comes great responsibility — especially when:

- You're deploying apps in the cloud
- There are many users, teams, and microservices
- A single misconfigured Pod or secret can compromise your system

So let's break down the main security areas.

■ 1. Access Control (Who can access what)

Use RBAC (Role-Based Access Control)

- Give users only the permissions they need.
- Example: Devs can view Pods, but only Admins can delete Deployments.

Use IRSA (IAM Roles for Service Accounts)

- Let Pods access AWS securely without hardcoded keys.
- Example: A logging pod gets permission to write to S3, and nothing else.
- Analogy: Think of RBAC like keycards in an office each person can only enter specific rooms.

2. Network Security

Use Network Policies to Control Pod-to-Pod Traffic

- Default in Kubernetes: All Pods can talk to each other
- Use Calico or Cilium to define who can talk to whom.

Isolate Environments

- Use **namespaces** to separate dev/test/prod.
- Use **Security Groups** to restrict access at the EC2/network level.
- Analogy: Think of Pods as rooms in a building network policies act like doors that can be locked or opened only for specific rooms.

3. Secrets Management

Prefer AWS Secrets Manager or Parameter Store

• These encrypt secrets with KMS and offer rotation, auditing, and fine-grained control.

Avoid:

- Plain YAML secrets
- Storing secrets directly in environment variables
- Tip: Use IRSA + AWS SDK or Secrets Store CSI Driver to fetch secrets at runtime.

4. Node Security

Use Managed Node Groups

• AWS patches the OS and handles lifecycle updates.

Choose Minimal AMIs:

- Bottlerocket (by AWS) or EKS-Optimized Amazon Linux 2
- Fewer packages = smaller attack surface

Apply regular updates

• Use **EKS Managed Node Groups** with scheduled updates or lifecycle automation.

5. Pod Security

- ▼ Enforce Pod Security Standards (PSS)
 - · Label namespaces with enforce=restricted
 - Prevent running as root, limit privilege escalation

Harden Pod Config:

- Use readOnlyRootFilesystem: true
- Drop Linux capabilities
- Add resource limits (CPU/memory)

Example:

yaml

securityContext:

runAsNonRoot: true

readOnlyRootFilesystem: true

capabilities: drop: ["ALL"]

6. Observability & Auditing

Enable:

- CloudTrail: Logs AWS API calls
- VPC Flow Logs: Logs network traffic
- EKS Audit Logs: Enable via CloudWatch for visibility into cluster events
- . GuardDuty: Detects unusual behavior or attacks
- Tip: Always monitor your cluster and set up alerts.

№ 7. Image and Supply Chain Security

- ✓ Use Trusted and Minimal Base Images
 - Avoid "ubuntu:latest" use slim or distroless images

Scan Images Before Deployment

- Tools: Trivy, Clair, Grype
- Can be automated in CI/CD
- Use Image Signing with Sigstore / Cosign
 - Verify that images are trusted and not tampered with

☑ Quick Recap: EKS Security Checklist

Area Best Practice
Access Control RBAC + IRSA

Network Security NetworkPolicies + Namespaces + Security Groups

Secrets Use AWS Secrets Manager + IRSA + CSI Driver

Nodes Use Bottlerocket / Amazon Linux 2 AMIs + auto patching

Pods Enforce PSS, no root, drop caps, read-only FS

Auditing Enable CloudTrail, VPC Flow Logs, Audit Logs

Area Best Practice

Image Security Scan, sign, and use minimal trusted images

How to Start

- 1. Set up RBAC (user access)
- 2. Enable IRSA for apps that use AWS services
- 3. Use Kubernetes Secrets initially, then migrate to AWS Secrets Manager
- 4. Use EKS Managed Node Groups
- 5. Label namespaces with restricted for pod security
- 6. Enable audit logging in EKS settings
- 7. Scan your container images using Trivy locally or in CI/CD

1. What is the purpose of using RBAC in EKS?

- A. To encrypt secrets inside Kubernetes
- B. To define which Pods can access external services
- C. To control which users or roles can access Kubernetes resources
- D. To control network traffic between Pods
- ✓ Answer: C

2. Why should you use IAM Roles for Service Accounts (IRSA) in EKS?

- A. To allow Secrets Manager to authenticate with Kubernetes
- B. To let pods securely access AWS services without hardcoding credentials
- C. To allow Kubelet to install pods faster
- D. To enable external CI/CD pipelines to deploy applications
- Answer: B

Where are Kubernetes Secrets stored by default in EKS?

- A. In S3 buckets
- B. In AWS Secrets Manager
- C. In DynamoDB
- D. In etcd (Kubernetes data store)
- Answer: D

What is the main limitation of Kubernetes Secrets by default?

- A. They cannot be accessed by pods
- B. They are encrypted with KMS
- C. They are stored as plaintext unless etcd encryption is enabled
- D. They require IRSA to be accessed
- ✓ Answer: C