



## EKS Deployment Flow

-

### STEP 1: Create Flask App

-----

#### Files:

- app.py
- requirements.txt

#### Code:

```
from flask import Flask
app = Flask(__name__)
@app.route("/")
def hello():
    return "Hello World from EKS"
```

#### Command:

```
> python app.py
```

Purpose: Have a working app before deployment.

-

-

### STEP 2: Containerize the App

-----

#### File: Dockerfile

#### Dockerfile:

```
FROM python:3.9-slim
WORKDIR /app
COPY requirements.txt .
RUN pip install -r requirements.txt
COPY app.py .
CMD ["python", "app.py"]
```

#### Commands:

```
> docker build -t flask-hello-world .
> docker run -p 5000:5000 flask-hello-world
```

**Purpose:** Package app into a container.

-

### STEP 3: Push to Amazon ECR

-----

#### Commands:

```
> aws ecr create-repository --repository-name flask-hello-world
> aws ecr get-login-password | docker login ...
> docker tag flask-hello-world <ecr-uri>
> docker push <ecr-uri>
```

**Purpose:** Store image on AWS for use in EKS.

yaml

-

### ⚙️ STEP 4: Create EKS Cluster

-----

**File:** cluster-config.yaml

#### YAML:

```
kind: ClusterConfig
metadata:
  name: flask-demo-cluster
  region: us-west-2
nodeGroups:
  - name: worker-nodes
    instanceType: t3.medium
```

#### Command:

```
> eksctl create cluster -f cluster-config.yaml
```

**Purpose:** Provision Kubernetes control plane + nodes.

-

-

### 🔗 STEP 5: Add Storage with PVC

-----

Files: storage-class.yaml, pvc.yaml

```
storage-class.yaml:
kind: StorageClass
provisioner: ebs.csi.aws.com
```

```
pvc.yaml:
kind: PersistentVolumeClaim
spec:
  resources:
    requests:
      storage: 1Gi
```

### Commands:

```
> kubectl apply -f storage-class.yaml
> kubectl apply -f pvc.yaml
```

Purpose: Provide durable disk space to your app.  
yaml

-

### STEP 6: Deploy Flask App

-----

File: flask-deployment.yaml

YAML (simplified):

kind: Deployment

spec:

replicas: 2

containers:

- name: flask-app

image: <your-ecr-image>

volumeMounts:

- mountPath: /app/data

name: flask-storage

volumes:

- name: flask-storage

persistentVolumeClaim:

claimName: flask-storage

Command:

```
> kubectl apply -f flask-deployment.yaml
```

Purpose: Start app pods in EKS using container image.  
yaml

-

### STEP 7: Create Internal Service

-----

File: flask-service.yaml

YAML:

kind: Service

spec:

selector:

app: flask-app

ports:

- port: 80

targetPort: 5000

type: ClusterIP

Command:

```
> kubectl apply -f flask-service.yaml
```

Purpose: Access pods inside the cluster.  
yaml

-

## STEP 8: Expose App to Internet

-----  
File: flask-ingress.yaml

YAML:

```
kind: Ingress
metadata:
  annotations:
    alb.ingress.kubernetes.io/scheme: internet-facing
spec:
  rules:
    - http:
        paths:
          - path: /
            backend:
              service:
                name: flask-service
                port:
                  number: 80
```

### Command:

```
> kubectl apply -f flask-ingress.yaml
> kubectl get ingress
```

Purpose: Give public access via AWS ALB.  
yaml

-

## STEP 9: Secure Access (RBAC + IAM)

-----  
File: rbac.yaml

YAML:

```
kind: Role
rules:
  - apiGroups: ["" ]
    resources: ["pods"]
    verbs: ["get", "list"]

kind: RoleBinding
subjects:
  - kind: ServiceAccount
    name: flask-sa
```

roleRef:  
  kind: Role  
  name: flask-role

Commands:  
> kubectl apply -f rbac.yaml

Purpose: Control what your app or users can access.  
yaml

-

## STEP 10: Harden Security in Deployment

-----

File: flask-deployment-secure.yaml

YAML:  
securityContext:  
  runAsNonRoot: true  
  runAsUser: 1000  
containers:  
  - securityContext:  
    allowPrivilegeEscalation: false  
    readOnlyRootFilesystem: true

Purpose: Prevent privilege abuse, follow security best practices.  
markdown

-

## STEP 11: Verify Everything

-----

Commands:  
> kubectl get nodes  
> kubectl get pods  
> kubectl logs <pod-name>  
> kubectl get svc  
> kubectl get ingress  
> curl <ALB-URL>

Purpose: Ensure everything is running and accessible.

---

## Summary

Step	Purpose	Key File / Command
1	Write app	app.py
2	Containerize	Dockerfile, docker build
3	Push to ECR	docker push, aws ecr
4	Create EKS Cluster	cluster-config.yaml, eksctl create
5	Add Storage	storage-class.yaml, pvc.yaml
6	Deploy app	flask-deployment.yaml

Step	Purpose	Key File / Command
7	Expose internally	flask-service.yaml
8	Expose externally (public)	flask-ingress.yaml
9	Add access control	rbac.yaml, kubectl apply
10	Harden container runtime	securityContext in deployment YAML
11	Test and verify	kubectl, curl, browser

## 1. How do you implement RBAC (Role-Based Access Control) in EKS?

**RBAC = "Who can do what?"**

RBAC stands for **Role-Based Access Control**, a system used in **Kubernetes** (and therefore in EKS too) to control **who can access what** resources in the cluster.

For example:

- **Dev users** can only view Pods.
- **Admins** can create or delete resources.
- **CI/CD pipelines** may get special access to certain namespaces.

---

### EKS + RBAC = IAM + Kubernetes RBAC

In **Amazon EKS**, you manage access at **two levels**:

1. **IAM level**: Controls which **users/roles from AWS** can authenticate into the cluster.
2. **Kubernetes RBAC level**: Controls **what those users can do** after logging in.

So even if an IAM user is allowed to log in to the cluster, RBAC decides what they can actually do inside Kubernetes.

---

### How IAM and RBAC Work Together in EKS

In EKS:

Layer	Controls	Example
<b>IAM (AWS)</b>	Who can authenticate to the cluster	IAM user kube_test_user
<b>RBAC (K8s)</b>	What that user can do inside Kubernetes	Role allowing list/get pods

✓ **IAM = Authentication**

✓ **RBAC = Authorization**

---

### Example: Grant kube\_test\_user read-only access to pods

---

#### 1 IAM User Setup (Assumed Already Exists)

You already have this IAM user:

ruby

-----

Name: kube\_test\_user

ARN: arn:aws:iam::<account-id>:user/kube\_test\_user

✓ Make sure the IAM user has permission to call eks:DescribeCluster and sts:AssumeRole.

---

## 2 Update aws-auth ConfigMap to Allow IAM Access

You must tell EKS:

“This IAM user maps to a Kubernetes user”.

Run:

bash

-----

kubectrl edit configmap aws-auth -n kube-system

Add under mapUsers::

yaml

-----

mapUsers:

- userarn: arn:aws:iam::<account-id>:user/kube\_test\_user
- username: kube\_test\_user
- groups:
- read-only-group

This says: "Let kube\_test\_user authenticate, and map them to the Kubernetes group read-only-group."

---

## 3 Create an RBAC Role for the Group

Create file: rbac-readonly.yaml

yaml

-----

apiVersion: rbac.authorization.k8s.io/v1

kind: Role

metadata:

- name: readonly-pods
- namespace: demo

rules:

- apiGroups: [""]
- resources: ["pods"]
- verbs: ["get", "list", "watch"]

---

apiVersion: rbac.authorization.k8s.io/v1

kind: RoleBinding

metadata:

- name: bind-readonly-pods
- namespace: demo

subjects:

- kind: Group
- name: read-only-group   # This matches the group in aws-auth
- apiGroup: rbac.authorization.k8s.io

roleRef:

- kind: Role
- name: readonly-pods

apiGroup: rbac.authorization.k8s.io

Apply it:

bash

-----

kubectl apply -f rbac-readonly.yaml

---

#### 4 Test Access (From kube\_test\_user CLI)

From a system where the kube\_test\_user is configured in AWS CLI:

bash

-----

aws eks update-kubeconfig --name your-cluster-name --region your-region \\  
--role-arn arn:aws:iam::<account-id>:user/kube\_test\_user

Try this:

bash

-----

kubectl get pods -n demo

✅ This should work.

Try this:

bash

-----

kubectl delete pod mypod -n demo

❌ This should fail – because the role is read-only.

---

#### 💡 Real-World Use Case

- You want **developers to view Pods** but **not delete anything**.
- You want **DevOps engineers** to manage deployments only in dev namespace.
- You want **CI/CD systems** to have scoped permissions, not full admin rights.

---

## 2. What is AWS IAM integration with EKS?

**Amazon EKS** (Elastic Kubernetes Service) runs **Kubernetes clusters** on AWS.

To secure it, we need to:

1. **Control who can access the cluster** (admins, devs, CI/CD users).
2. **Control what running Pods can access in AWS** (like S3, DynamoDB, etc.).

To do this, EKS integrates with **AWS IAM (Identity and Access Management)**.

---

### Two Key Parts of IAM Integration in EKS

Purpose	Technique	Use Case
<b>Cluster access (users)</b>	aws-auth ConfigMap	Map IAM users/roles to Kubernetes users
<b>Pod access to AWS services</b>	<b>IAM Roles for Service Accounts (IRSA)</b>	Let apps in pods access S3, DynamoDB, etc. securely

---



## ✅ Part 1: IAM → Kubernetes Access (via aws-auth ConfigMap)

### 📌 What it does:

Maps **IAM Users or Roles** to **Kubernetes RBAC identities**.

### 📁 Where?

In a special ConfigMap inside the EKS cluster:

```
bash
```

```
-----
```

```
kubectl edit configmap aws-auth -n kube-system
```

### 📄 Example:

```
yaml
```

```
-----
```

```
mapUsers:
```

```
- userarn: arn:aws:iam::123456789012:user/dev
```

```
  username: dev-user
```

```
  groups:
```

```
  - system:masters
```

This allows:

- IAM user dev to connect to the EKS cluster as dev-user
- Grants them full access (system:masters) unless you restrict with RBAC

### 🔑 Check access:

```
bash
```

```
-----
```

```
kubectl auth can-i list pods --as dev-user
```

---

## ✅ Part 2: IAM Roles for Service Accounts (IRSA)

### 🔗 Problem:

Let's say a **pod** needs to access an **S3 bucket**. How do we give it access **safely**, without hardcoding AWS credentials?

### ✅ Solution: Use IRSA (IAM Roles for Service Accounts)

**IRSA = Secure way to attach an IAM Role to a Kubernetes ServiceAccount.**

So, the Pod assumes that IAM Role, and gets permissions to AWS services — no hardcoded creds.

---

### 🔄 How IRSA Works

1. ✅ You create an **IAM Role** with the needed AWS permissions.
2. ✅ Attach an **OIDC trust policy**, so EKS knows this role can be assumed by a Pod.
3. ✅ In Kubernetes, you create a **ServiceAccount** linked to that IAM Role.
4. ✅ Your **Pod uses that ServiceAccount**, and it gets AWS access automatically.

---

### 📄 Example YAML for IRSA:

**IAM Role trust policy:**

```
json
```

```
-----
```

```
{
```

```
  "Effect": "Allow",
```

```
  "Principal": {
```

```

    "Federated": "arn:aws:iam::123456789012:oidc-provider/oidc.eks.us-west-2.amazonaws.com/id/EXAMPLED539D4633E53DE1B716D3041E"
  },
  "Action": "sts:AssumeRoleWithWebIdentity",
  "Condition": {
    "StringEquals": {
      "oidc.eks...:sub": "system:serviceaccount:default:myapp-sa"
    }
  }
}

```

### Kubernetes ServiceAccount:

yaml

-----

apiVersion: v1

kind: ServiceAccount

metadata:

name: myapp-sa

annotations:

eks.amazonaws.com/role-arn: arn:aws:iam::123456789012:role/MyAppIRSA

Then in your Deployment/Pod spec:

yaml

-----

serviceAccountName: myapp-sa

Now the pod can access AWS services like S3 or DynamoDB with **fine-grained permissions**.

---

### Summary for Beginners

Feature	Purpose	Where It Works
aws-auth ConfigMap	Maps <b>IAM users/roles</b> to <b>K8s users</b>	Controls <b>who</b> can access the EKS cluster
<b>IRSA</b>	Maps <b>IAM roles</b> to <b>Pods</b>	Controls what <b>apps</b> inside EKS can access in AWS

---

### Real-Life Use Cases

- A **developer** logs in using their IAM user to view resources: aws-auth mapping.
  - A **Pod uploads files to S3**: IRSA.
  - A **CI/CD pipeline** (e.g., GitHub Actions) assumes an IAM role and deploys to EKS.
- 

---

## 3. How do you configure service accounts in EKS?

### a First: What's the Problem We're Solving?

Imagine you have a **Pod in EKS** that needs to access:

- An **S3 bucket** to upload logs
- Or a **DynamoDB table** to read/write data

But — **how does that Pod get AWS credentials securely?**

We don't want to:

- Hardcode AWS keys in the Pod
- Use overly broad permissions

✅ **Solution: Use IAM Roles for Service Accounts (IRSA)**

With IRSA, you attach a specific **IAM Role** to a **Kubernetes ServiceAccount**.

Pods using that ServiceAccount will **automatically get AWS credentials** — securely and scoped.

---

🔧 **Steps to Configure IRSA in EKS (Beginner Friendly)**

---

◆ **Step 1: Enable IAM OIDC Provider**

Each EKS cluster needs an **OIDC provider** (OpenID Connect) so IAM can trust Kubernetes.

✅ Run this command:

bash

-----

```
eksctl utils associate-iam-oidc-provider \
  --region <your-region> \
  --cluster <your-cluster-name> \
  --approve
```

This allows your IAM roles to trust tokens from your EKS cluster.

---

◆ **Step 2: Create IAM Policy**

Create a policy with **least-privilege AWS permissions**.

✅ Example: S3 Read/Write

json

-----

```
{
  "Version": "2012-10-17",
  "Statement": [{
    "Effect": "Allow",
    "Action": [
      "s3:ListBucket",
      "s3:GetObject",
      "s3:PutObject"
    ],
    "Resource": [
      "arn:aws:s3:::my-bucket",
      "arn:aws:s3:::my-bucket/*"
    ]
  }]
}
```

Save it as s3-access-policy.json, then run:

bash

```
-----  
aws iam create-policy \  
  --policy-name S3AccessPolicy \  
  --policy-document file://s3-access-policy.json
```

---

### ◆ Step 3: Create IAM Role with Trust Policy

This IAM Role:

- Can only be assumed by EKS Pods
- Only when using a specific ServiceAccount

✓ Use this trust policy (replace placeholders):  
json

```
-----  
{  
  "Version": "2012-10-17",  
  "Statement": [{  
    "Effect": "Allow",  
    "Principal": {  
      "Federated": "arn:aws:iam::<account-id>:oidc-  
provider/oidc.eks.<region>.amazonaws.com/id/<eks-id>"  
    },  
    "Action": "sts:AssumeRoleWithWebIdentity",  
    "Condition": {  
      "StringEquals": {  
        "oidc.eks.<region>.amazonaws.com/id/<eks-id>:sub":  
"system:serviceaccount:default:s3-access-sa"  
      }  
    }  
  }  
}]  
}
```

Then create the IAM role:

bash

```
-----  
aws iam create-role \  
  --role-name S3AccessRole \  
  --assume-role-policy-document file://trust-policy.json
```

And attach the S3 policy to it:

bash

```
-----  
aws iam attach-role-policy \  
  --role-name S3AccessRole \  
  --policy-arn arn:aws:iam::<account-id>:policy/S3AccessPolicy
```

---

### ◆ Step 4: Create a Kubernetes ServiceAccount (with annotation)

This is the **link between your pod and the IAM role**.

📄 Example YAML:

yaml

```
-----
```

```
apiVersion: v1
kind: ServiceAccount
metadata:
  name: s3-access-sa
  annotations:
    eks.amazonaws.com/role-arn: arn:aws:iam::<account-id>:role/S3AccessRole
```

Apply it:

bash

-----

kubectl apply -f sa.yaml

---

### ◆ Step 5: Use the ServiceAccount in Your Pod/Deployment

Now update your Pod or Deployment YAML to **use the service account**.

📄 Example:

yaml

-----

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: myapp
spec:
  replicas: 1
  selector:
    matchLabels:
      app: myapp
  template:
    metadata:
      labels:
        app: myapp
    spec:
      serviceAccountName: s3-access-sa
      containers:
        - name: mycontainer
          image: myapp:latest
          # App can now access AWS S3 via SDK without hardcoding credentials
```

---

### ✅ Why Is This Great?

- 🛡️ **Secure:** No need to store access keys in pods
- 🎯 **Fine-grained access:** Each pod gets only the permissions it needs
- 📦 **Cloud-native:** Works seamlessly with AWS SDKs and CLI in your app

---

### Summary

#### Step What You Do

- 1 Enable OIDC
- 2 Create IAM Policy
- 3 Create IAM Role with Trust

#### Why

So IAM can trust EKS tokens  
Define what AWS resources are allowed  
Link the IAM Role to EKS ServiceAccount

## Step What You Do

## Why

- |                                               |                                               |
|-----------------------------------------------|-----------------------------------------------|
| <div>4</div> Create Kubernetes ServiceAccount | Annotate it with IAM Role ARN                 |
| <div>5</div> Use ServiceAccount in Pod        | Your app gets secure AWS access automatically |
- 

---

## 4. What are Pod Security Standards and how do you implement them?



### In simple terms:

PSS are **rules that control how secure your Pods are** — like what they're allowed to do, and what they're **not** allowed to do.

Kubernetes defines **3 security levels**:

Level	Purpose
-------	---------

<b>privileged</b>	No restrictions. Good for debugging, but unsafe for production.
-------------------	-----------------------------------------------------------------

<b>baseline</b>	Prevents known bad practices (e.g., running as root). Safer for most workloads.
-----------------	---------------------------------------------------------------------------------

<b>restricted</b>	Very strict. Follows best practices for least privilege and hardening. Ideal for production.
-------------------	----------------------------------------------------------------------------------------------

---

### ! Why Do You Need PSS in EKS?

Without any restrictions:

- Pods could run as **root**
- Mount **host paths** (accessing the EC2 machine itself!)
- Use **privileged mode** (full admin rights inside container)

This opens up security risks like container breakout attacks.

So, PSS lets you **enforce rules** at the **namespace level** to protect the cluster.



### How to Implement Pod Security Standards (Step-by-Step)

Kubernetes 1.25+ (used in modern EKS versions) comes with a built-in **PodSecurity admission controller**.



#### Step 1: Choose Your Namespace

Let's say your namespace is dev.

You want all pods in dev to follow restricted security policies.



#### Step 2: Add Labels to Enforce PSS

Use this command:

bash

-----

```
kubectl label namespace dev \
  pod-security.kubernetes.io/enforce=restricted
```



This means: **Pods in the dev namespace must follow the restricted policy.**

You can also add optional labels for **audit** or **warn**:

bash

-----

pod-security.kubernetes.io/warn=baseline

pod-security.kubernetes.io/audit=privileged

So Kubernetes will **warn** or **log** if a pod violates less strict policies.



### Example Label Setup:

Label	Meaning
pod-security.kubernetes.io/enforce=restricted	Block anything below restricted
pod-security.kubernetes.io/warn=baseline	Show warning if policy is below baseline
pod-security.kubernetes.io/audit=privileged	Log policy violations of privileged level



### Example: What Gets Blocked?

If someone tries to apply this Pod YAML:

yaml

-----

apiVersion: v1

kind: Pod

metadata:

  name: badpod

spec:

  containers:

    - name: alpine

      image: alpine

      securityContext:

        privileged: true

It will get **denied** in a namespace labeled restricted because:

- Running a privileged container is **not allowed** under restricted mode.



### Bonus: Advanced Policy with OPA Gatekeeper or Kyverno

If you want **custom rules** (e.g., "pods must use a specific image registry" or "no containers with hostPath"), use:

Tool	Use Case
<b>OPA Gatekeeper</b>	Policy as code using Rego language
<b>Kyverno</b>	Simpler YAML-based policy engine, very Kubernetes-native

These tools provide **richer validation**, auditing, and auto-remediation.



### Summary for Beginners

Concept	Meaning
<b>Pod Security Standards (PSS)</b>	Built-in security levels for Pods
<b>privileged</b>	No restrictions (not safe)
<b>baseline</b>	Medium safety — blocks known bad configs
<b>restricted</b>	Strong safety — best for production
<b>Namespace labels</b>	Used to enforce security levels

Concept	Meaning
OPA/Gatekeeper/Kyverno	Used for custom advanced policy enforcement

### 💡 Real-World Use Case in EKS

- You run a **multi-tenant EKS cluster** (e.g., dev, staging, prod).
- You want to make sure **developers can't run insecure containers**.
- You label dev and prod namespaces with different PSS levels.
- In prod, only hardened, secure workloads can run.

## 5. How do you handle secrets management in EKS?

### s🔒 What Are Secrets and Why Manage Them?

In applications, **secrets** = sensitive info like:

- Database passwords
- API tokens
- SSH keys
- AWS credentials

If these secrets are not stored securely, **anyone who compromises a pod** can steal them — leading to data leaks or system compromise.

So in EKS (Kubernetes on AWS), managing secrets safely is **critical**.

### 🔒 3 Options for Secrets Management in EKS

Let's go through them **from basic to advanced**, so beginners can understand and build confidence.

#### ✅ Option 1: Kubernetes Secrets (Basic)

- Secrets are stored inside Kubernetes using built-in Secret resources.
- They are **base64-encoded**, not encrypted by default.
- Stored in **etcd**, which should be encrypted (in EKS, it is by default).

#### 📄 Example YAML:

```
yaml
-----
apiVersion: v1
kind: Secret
metadata:
  name: db-creds
type: Opaque
data:
  username: YWRtaW4=    # "admin"
  password: MWYyZDFlMmU2N2Rm # "1f2d1e2e67df"
```

To create it:

```
bash
```

```
-----
```

```
kubectl apply -f secret.yaml
```

Then in a Pod:

```
yaml
```

```
-----
```



env:

```
- name: DB_USER
  valueFrom:
    secretKeyRef:
      name: db-creds
      key: username
```

#### ⚠ **Limitations:**

- Not encrypted at rest (unless etcd encryption is configured).
- Secrets visible to anyone with access to the namespace.
- Better for **non-critical** dev/test use.

---

### ✅ **Option 2: AWS Secrets Manager with IRSA (Production-Ready)**

AWS Secrets Manager is a **fully managed secrets service**:

- Secrets are encrypted with KMS.
- You can **rotate secrets automatically**.
- Access is controlled via **IAM** (and in EKS, via **IRSA**).

#### 🔄 **Flow:**

1. Store your secret in AWS Secrets Manager.
2. Create a Kubernetes **ServiceAccount with IRSA**.
3. Pod uses AWS SDK (or init/sidecar container) to **fetch the secret securely**.

#### 🔑 **Example:**

Store a secret:

bash

-----

```
aws secretsmanager create-secret \
  --name db-creds \
  --secret-string '{"username":"admin","password":"s3cr3t"}
```

In your app (Node.js, Python, etc.), use the AWS SDK:

js

-----

```
const client = new SecretsManager();
const secret = await client.getSecretValue({ SecretId: 'db-creds' }).promise();
```

✅ Best for production. Supports **audit logging, rotation, and IAM-based fine control**.

---

### ✅ **Option 3: Secrets Store CSI Driver (Recommended Integration)**

This lets Kubernetes **mount external secrets** (like from Secrets Manager or Vault) as **volumes**, so your app doesn't need to call AWS SDK directly.

This is great if:

- You don't want to modify your app code.
- You want secret values injected directly as files.

#### 🔑 **How It Works:**

1. Install **Secrets Store CSI Driver** in your EKS cluster.
2. Install AWS provider plugin.
3. Create a SecretProviderClass that tells which secrets to pull.
4. Mount those secrets in Pod using a CSI volume.

#### 📄 **Example:**

yaml

```
-----
apiVersion: secrets-store.csi.x-k8s.io/v1
kind: SecretProviderClass
metadata:
  name: aws-secrets
spec:
  provider: aws
  parameters:
    objects: |
      - objectName: "db-creds"
        objectType: "secretsmanager"
```

Then in your pod:

yaml

```
-----
volumeMounts:
- name: secrets-store
  mountPath: "/mnt/secrets"
```

```
volumes:
- name: secrets-store
  csi:
    driver: secrets-store.csi.k8s.io
    readOnly: true
    volumeAttributes:
      secretProviderClass: "aws-secrets"
```

Now /mnt/secrets/db-creds contains the secret securely — no need for base64 decoding or SDK calls.

---

### Summary for Beginners

Method	Description	Best For
<b>Kubernetes Secrets</b>	Built-in, simple, stored in etcd	Dev/test environments
<b>AWS Secrets Manager + IRSA</b>	Securely access AWS secrets via IAM roles	Production apps needing AWS secrets
<b>Secrets Store CSI Driver</b>	Mount secrets as volumes from AWS Secrets Manager/Vault	No-code change, production-grade

---

### Recommendation for Beginners in EKS

- **Start with Kubernetes Secrets** in test/dev.
- Learn **IRSA** to fetch from **AWS Secrets Manager** in production.
- Use **Secrets Store CSI Driver** for **secure volume-based access** — especially useful in regulated environments.

---

## 6. What are the security best practices for EKS?

### Overview: Why Security in EKS Matters

EKS (Elastic Kubernetes Service) gives you the power of Kubernetes on AWS.

But with great power comes great responsibility — especially when:

- You're deploying apps in the cloud
- There are many users, teams, and microservices
- A single misconfigured Pod or secret can compromise your system

So let's break down the **main security areas**.



## 1. Access Control (Who can access what)



### Use RBAC (Role-Based Access Control)

- Give users only the **permissions they need**.
- Example: Devs can view Pods, but only Admins can delete Deployments.



### Use IRSA (IAM Roles for Service Accounts)

- Let Pods access AWS securely without hardcoded keys.
- Example: A logging pod gets permission to write to S3, and nothing else.



Analogy: Think of RBAC like keycards in an office — each person can only enter specific rooms.



## 2. Network Security



### Use Network Policies to Control Pod-to-Pod Traffic

- Default in Kubernetes: All Pods can talk to each other 🤖
- Use Calico or Cilium to define **who can talk to whom**.



### Isolate Environments

- Use **namespaces** to separate dev/test/prod.
- Use **Security Groups** to restrict access at the EC2/network level.



Analogy: Think of Pods as rooms in a building — network policies act like doors that can be locked or opened only for specific rooms.



## 3. Secrets Management



### Prefer AWS Secrets Manager or Parameter Store

- These encrypt secrets with KMS and offer rotation, auditing, and fine-grained control.



### Avoid:

- Plain YAML secrets
- Storing secrets directly in environment variables



Tip: Use IRSA + AWS SDK or Secrets Store CSI Driver to fetch secrets at runtime.



## 4. Node Security



### Use Managed Node Groups

- AWS patches the OS and handles lifecycle updates.



### Choose Minimal AMIs:

- **Bottlerocket** (by AWS) or **EKS-Optimized Amazon Linux 2**
- Fewer packages = smaller attack surface



### Apply regular updates

- Use **EKS Managed Node Groups** with scheduled updates or lifecycle automation.



## 5. Pod Security

### ✅ Enforce Pod Security Standards (PSS)

- Label namespaces with enforce=restricted
- Prevent running as root, limit privilege escalation

### ✅ Harden Pod Config:

- Use readOnlyRootFilesystem: true
- Drop Linux capabilities
- Add resource limits (CPU/memory)

🧠 Example:

yaml

-----

```
securityContext:
  runAsNonRoot: true
  readOnlyRootFilesystem: true
  capabilities:
    drop: ["ALL"]
```

---

## 👁️ 6. Observability & Auditing

### ✅ Enable:

- **CloudTrail:** Logs AWS API calls
- **VPC Flow Logs:** Logs network traffic
- **EKS Audit Logs:** Enable via CloudWatch for visibility into cluster events
- **GuardDuty:** Detects unusual behavior or attacks

🧠 Tip: Always **monitor your cluster and set up alerts**.

---

## 📦 7. Image and Supply Chain Security

### ✅ Use Trusted and Minimal Base Images

- Avoid "ubuntu:latest" — use slim or distroless images

### ✅ Scan Images Before Deployment

- Tools: **Trivy, Clair, Gype**
- Can be automated in CI/CD

### ✅ Use Image Signing with Sigstore / Cosign

- Verify that images are trusted and not tampered with

---

### ✅ Quick Recap: EKS Security Checklist

Area	Best Practice
Access Control	RBAC + IRSA
Network Security	NetworkPolicies + Namespaces + Security Groups
Secrets	Use AWS Secrets Manager + IRSA + CSI Driver
Nodes	Use Bottlerocket / Amazon Linux 2 AMIs + auto patching
Pods	Enforce PSS, no root, drop caps, read-only FS
Auditing	Enable CloudTrail, VPC Flow Logs, Audit Logs
Image Security	Scan, sign, and use minimal trusted images

---

### How to Start

1. **Set up RBAC** (user access)

2. **Enable IRSA** for apps that use AWS services
3. **Use Kubernetes Secrets initially**, then migrate to **AWS Secrets Manager**
4. **Use EKS Managed Node Groups**
5. **Label namespaces with restricted** for pod security
6. **Enable audit logging** in EKS settings
7. **Scan your container images** using Trivy locally or in CI/CD

---

### 1. What is the purpose of using RBAC in EKS?

- A. To encrypt secrets inside Kubernetes
- B. To define which Pods can access external services
- C. To control which users or roles can access Kubernetes resources
- D. To control network traffic between Pods

 **Answer: C**

---

### 2. Why should you use IAM Roles for Service Accounts (IRSA) in EKS?

- A. To allow Secrets Manager to authenticate with Kubernetes
- B. To let pods securely access AWS services without hardcoding credentials
- C. To allow Kubelet to install pods faster
- D. To enable external CI/CD pipelines to deploy applications


 **Answer: B**

---

---

### Where are Kubernetes Secrets stored by default in EKS?

- A. In S3 buckets
- B. In AWS Secrets Manager
- C. In DynamoDB
- D. In etcd (Kubernetes data store)

 **Answer: D**

---

### What is the main limitation of Kubernetes Secrets by default?

- A. They cannot be accessed by pods
- B. They are encrypted with KMS
- C. They are stored as plaintext unless etcd encryption is enabled
- D. They require IRSA to be accessed

 **Answer: C**

---