



## Scenario based questions : EKS

1. Your team needs to deploy a microservices-based application on AWS using Kubernetes. How would you set up the EKS cluster?

### Steps

#### 1. VPC Setup

- Use an existing VPC or create a new one with:
  - **Public and private subnets** across multiple Availability Zones (for high availability).
  - Enable **internet gateway** and **NAT gateways** for outbound access.

#### 2. Create EKS Cluster

- Go to **EKS Console** or use **eksctl** (simpler CLI tool).
- Example eksctl command:

```
bash
```

```
-----
```

```
eksctl create cluster \  
  --name my-eks-cluster \  
  --region us-east-1 \  
  --nodes 2 \  
  --node-type t3.medium \  
  --managed
```

#### 3. Configure kubectl

- Update the kubectl config to talk to the new cluster:

```
bash
```

```
-----
```

```
aws eks --region us-east-1 update-kubeconfig --name my-eks-cluster
```

#### 4. Deploy Core Add-ons

- Install essential Kubernetes add-ons:
  - **VPC CNI plugin** (networking)
  - **CoreDNS** (service discovery)
  - **kube-proxy** (network routing)
  - Optionally: **AWS Load Balancer Controller**, **Cluster Autoscaler**, **Metrics Server**

#### 5. Secure the Cluster

- Use **IAM roles for service accounts (IRSA)** to allow pods access to AWS services securely.

- Apply **RBAC** to restrict user and application permissions.
- Enable **logging** using CloudWatch.

## 6. Deploy Microservices

- Use `kubectl apply -f <deployment.yaml>` or **Helm** to deploy your microservices.
- Use **Kubernetes Services** (ClusterIP, NodePort, or LoadBalancer) to expose them.
- Organize using **Namespaces** (e.g., dev, test, prod) for logical separation.

## 7. Monitor & Auto-scale

- Enable **Horizontal Pod Autoscaler (HPA)** and **Cluster Autoscaler**.
- Use **CloudWatch Container Insights** or **Prometheus + Grafana** for observability.

## 8. CI/CD Integration (Optional)

- Set up pipelines using **CodePipeline**, **GitHub Actions**, or **Jenkins** to automate deployments.

## 2. Scenario:

You have a backend app and a frontend React app that need to be deployed on EKS. Walk us through the CI/CD deployment strategy.

### CI/CD Deployment Strategy to EKS

#### 1. Prepare Your Code Repositories

- **Frontend** and **Backend** live in separate GitHub repositories (or in separate folders of a mono-repo).
- Each project contains:
  - Dockerfile (for containerizing)
  - deployment.yaml and service.yaml files (for Kubernetes deployment)

---

#### 2. Build & Push Docker Images (CI)

For both apps, I would:

- Set up CI using **GitHub Actions**, **CodeBuild**, or **Jenkins**.

#### CI Workflow Steps:

1. **Checkout Code**
2. **Build Docker Image**
3. **Tag Image** with commit SHA or version
4. **Push Image to Amazon ECR**

#### Example CI for backend (GitHub Actions):

yaml

-----

- name: Build and Push Docker Image

run: |

docker build -t \$ECR\_URI/backend:\$GITHUB\_SHA .

docker push \$ECR\_URI/backend:\$GITHUB\_SHA

Repeat similarly for frontend app.

---

### 3. Kubernetes Manifests (with Image Tags)

- Use kustomize or helm or dynamic sed to update the image tags in the manifest.

Example:

yaml

-----

containers:

- name: backend

image: <ECR\_URI>/backend:<GITHUB\_SHA>

---

### 4. Deploy to EKS (CD)

- Use kubectl or helm to deploy updated manifests to the EKS cluster.

#### CD Workflow:

- Authenticate using:

bash

-----

aws eks update-kubeconfig --name my-cluster

- Apply manifests:

bash

-----

kubectl apply -f backend-deployment.yaml

kubectl apply -f frontend-deployment.yaml

---

### 5. Automate the Pipeline

Put everything into CI/CD stages:

Stage	Tool Used	Purpose
CI Build	GitHub Actions	Build and push Docker images to ECR
CD Deploy	GitHub Actions / ArgoCD / CodePipeline	Deploy latest changes to EKS

---

### 6. Expose Services

- Use LoadBalancer type service for frontend.
- Use ClusterIP for backend (accessed internally), or Ingress for routing.

---

### 7. Post-deploy Checks

- Run health checks and notify if deployment failed.
- Optionally: Rollback if required using Helm or ArgoCD.

---

### Why this strategy?

- Ensures **automation**, **repeatability**, and **version control**.
  - Can easily extend to multiple environments (dev, staging, prod).
  - Uses AWS-native services like **ECR**, **EKS**, and optionally **CodePipeline** or **ArgoCD**.
-

### 3. A service running in EKS can't connect to an RDS database in a private subnet. How would you debug this?

To debug a service in **EKS** that cannot connect to an **RDS** database in a **private subnet**, I would follow this **layered approach**:

---

#### Debugging Checklist



---

##### 1. Confirm RDS Endpoint & Port

- Double-check:
  - Correct **RDS hostname** (endpoint)
  - Correct **port** (usually 5432 for PostgreSQL, 3306 for MySQL)

---

##### 2. VPC Networking Check

-  Ensure **EKS worker nodes** (EC2) are in the **same VPC** as RDS, or **VPC peering** is in place.
-  Confirm RDS is in a **private subnet**, and the **Kubernetes pods** can access private subnets.
- Use `kubectl get nodes -o wide` to see EC2 node IPs and subnet info.

---

##### 3. Security Groups

- Check RDS **security group inbound rules**:
  - Allow traffic from EKS **worker node security groups** or specific CIDR.
  - For example:

yaml

-----

Type: TCP

Port: 5432

Source: sg-xxxx (EKS Node SG) or 10.0.0.0/16

---

##### 4. NACL (Network ACL) Rules

- Ensure **subnet-level NACLs** are not blocking traffic:
  - Both **inbound** and **outbound** rules should allow traffic on the correct port.

---

##### 5. DNS Resolution

- Run a test pod to check if RDS hostname resolves:

bash

-----

`kubectl run curlpod --image=amazonlinux --restart=Never -it -- bash`

`yum install -y bind-utils`

`nslookup mydb.xxxxxxxxxx.rds.amazonaws.com`

---

##### 6. Connectivity Test

- From the pod, test DB connectivity:

bash

-----

curlpod\$ telnet <rds-endpoint> 5432

Or use psql or mysql client to connect and confirm error messages.

---

## 7. IAM or DB Auth (Optional)

- If you're using **IAM authentication**:
  - Ensure proper IAM roles and token generation.
- Also check **DB-level user credentials** are correct.

---

## 8. Review Logs

- Check:
  - Pod logs: `kubectl logs <pod>`
  - App logs for connection timeouts or errors
  - RDS logs (from AWS Console) for rejected connections

---

## 9. Restart or Rollout Pod

- If you made config changes (like adding new env vars for DB endpoint), restart the pod:

bash

-----

`kubectl rollout restart deployment my-backend`

---

## 10. Add NAT Gateway (If Needed)

- If pods are in private subnets and need **outbound internet** (e.g., for DNS resolution), ensure there's a **NAT Gateway**.

---

### Bonus Tip:

- Use a **debug container** (`curl`, `netcat`, `busybox`) to test connectivity from within the cluster.

---


## Summary

This issue often boils down to **VPC misconfiguration**, **security groups**, or **missing DNS/NAT settings**. By validating **network path**, **security layers**, and **service discovery**, we can quickly isolate and fix the problem.

Which add-on is **mandatory** for Kubernetes networking in EKS?

- A. CoreDNS
- B. VPC CNI plugin
- C. Metrics Server
- D. Cluster Autoscaler

 **Answer: B**

 **Reason:** The VPC CNI plugin is required for pod networking in AWS VPC.

---

---

For the frontend React app exposed to the internet, which Kubernetes Service type is best?

- A. ClusterIP
- B. NodePort
- C. LoadBalancer
- D. Headless Service

✅ **Answer: C**

💡 **Reason:** LoadBalancer type automatically provisions an AWS ELB to expose services externally.

If DNS resolution fails inside the pod when connecting to RDS, which command helps check the hostname?

- A. ping
- B. nslookup
- C. netstat
- D. curl

✅ **Answer: B**

💡 **Reason:** nslookup verifies DNS resolution inside the cluster network.