


@devopschallengehub



1. Your EC2 CPU utilization crossed 80%. How would you set up a notification?

Answer:

I'd create a **CloudWatch Alarm** based on the EC2's CPU metric. Here's what I'd do:

- Go to CloudWatch → Alarms → Create Alarm
- Choose the **CPUUtilization metric** for that EC2 instance
- Set the threshold to **> 80% for 2 consecutive periods of 5 minutes** (or 1 minute if detailed monitoring is enabled)
- Then **attach an SNS topic** (Simple Notification Service)
- Subscribe an email/Slack/Lambda to the SNS topic

Example:

In one project, we had a backend API server under occasional load spikes. I set up a CloudWatch alarm that notified the on-call DevOps email when CPU crossed 80% — that helped us detect when autoscaling was needed.


If an EC2 instance has **basic monitoring** enabled, what is the minimum period you can set for the CPU alarm?

- A) 1 minute
- B) 5 minutes
- C) 30 seconds
- D) Depends on EC2 size

B) 5 minutes ✓

A CloudWatch alarm on **CPUUtilization > 80% for 2 consecutive 5-minute periods** will trigger after:

- A) Exactly 5 minutes of high CPU
- B) Exactly 10 minutes of high CPU
- C) 80% CPU at any moment
- D) Depends on CloudWatch evaluation delay

B) Exactly 10 minutes of high CPU 

(It's 2 periods × 5 minutes before state changes.)

Which statement about CloudWatch CPU alarms is **false**?

- A) The metric name is case-sensitive (CPUUtilization).
- B) An alarm can trigger an SNS notification.
- C) You must install the CloudWatch agent to get CPU metrics.
- D) Alarm states can be OK, ALARM, or INSUFFICIENT_DATA.

C) You must install the CloudWatch agent to get CPU metrics. 

You set a CPU alarm at **80%** for a burstable `t2.micro`. The CPU spikes to 90% for 30 seconds, then drops to 20%. What happens? (AWS CloudWatch)

- A) Alarm triggers immediately
- B) Alarm triggers after 2 periods
- C) Alarm never triggers
- D) Alarm triggers only if credit balance is low

C) Alarm never triggers 
