



How do you define cache behaviour in CloudFront to optimize performance?

Cache behavior in CloudFront defines **how requests are handled and cached**. You control it using **policies, TTLs, and cache keys**:

- ◆ **Cache Policies** – Decide what to cache (e.g., cache images but forward API calls).
- ◆ **TTL (Time-to-Live)** – Defines how long objects stay cached (short TTL for HTML, long TTL for images).
- ◆ **Cache Keys** – Parameters (headers, cookies, query strings) that CloudFront uses to uniquely identify cached objects.

Example:

- A React app: cache .js and .css for 24h, but API responses for only 60s (because API data changes fast).
- Use cache keys carefully: If you cache API results with query strings, each query becomes a separate cache entry.

DevOps Pipeline Use-case:

When deploying apps, pipeline scripts can update cache behaviors automatically (via Terraform/CloudFormation). For example, after pushing a new frontend release, pipeline triggers **cache invalidation + updates TTLs** so users get fresh code but images remain cached.

Which of the following best describes a Cache Policy in CloudFront?

- A) Defines how long CloudFront caches objects
- B) Controls which headers, cookies, and query strings are used as cache keys
- C) Decides what to cache and what to forward to the origin
- D) Determines where CloudFront serves traffic from

✓ **Answer: C**

Explanation: A Cache Policy tells CloudFront *what to cache and what to forward* (e.g.,

cache images but forward API calls). TTL is separate, and location is determined by Edge Locations.

In CloudFront, TTL (Time-to-Live) primarily controls:

- A) Which cookies are used in the cache key
- B) How long objects stay cached before revalidation
- C) Which origin is selected for requests
- D) Whether API responses are forwarded

✅ **Answer: B**

Explanation: TTL defines *how long objects are cached*. Example: HTML cached for 60s, images for 24h.

What happens if you include query strings in the cache key for API caching?

- A) All API responses are cached under the same key
- B) Each unique query string generates a separate cache entry
- C) CloudFront ignores the query string and caches only one version
- D) CloudFront invalidates all cache entries automatically

✅ **Answer: B**

Explanation: Cache keys use query strings to differentiate cache entries. ?id=1 and ?id=2 become separate cached objects.