

@devopschallengehub



## How would you perform a zero-downtime schema migration on RDS ?

### STAR

#### Situation

We had an **e-commerce app running on MySQL RDS**, and business wanted to add a new column `order_status` to support advanced tracking — things like “order packed”, “shipped”, “out for delivery”.

The problem was, the **orders table had millions of records**, and a simple `ALTER TABLE` would **lock the table** for several minutes — which meant customers couldn’t place or update orders during that time.

That was **unacceptable during live traffic hours**.

#### Task

My job was to **add this new column and index** to improve query speed — **without any downtime** or impact on customers.

So, I had to plan a **zero-downtime schema migration**.

#### Action

Here’s how I approached it step by step:

1. **Plan and test safely first:**

I took an **RDS snapshot** and restored it in a **staging environment**.

I ran the schema migration there to measure how long it might take and whether it caused any locks.

2. **Used an online schema migration tool:**

Instead of a direct `ALTER TABLE`, I used **gh-ost** — it’s an open-source tool by GitHub that performs schema changes **in the background** without blocking writes.

It creates a shadow table, copies data slowly, applies live changes, and then swaps tables almost instantly.

3. **Implemented backward compatibility:**

I ensured that the **new app code** could handle both old and new schemas.

So, for a while, the app was doing **dual writes** — writing both to old and new columns.

This meant we could safely test the new setup without breaking anything.

4. **Did a batch backfill:**

Since the new column was empty for existing rows, I ran a **background job** that filled old records in **small batches** during off-peak hours, while monitoring **CPU, IOPS, and replica lag** via CloudWatch.

5. **Monitored continuously:**

I kept an eye on metrics like **replication lag** and **query performance**, to ensure there was no load spike or replication issue.

6. **Switched gradually:**

Once all data was in sync and verified, we flipped the **feature flag** in the app so it started reading from the new column.

After a week of stability, we safely **dropped the old field**.

---

### Result

- ✓ The migration was completed with **zero downtime** — customers didn't notice anything.
- ✓ The new index improved query performance by **~40%**.
- ✓ Our approach became a **standard playbook** for all future schema changes in the company.

---

### Interview

“**backward/forward compatibility**, test in staging with production-like data, use **online schema tools** and replicas, implement **dual-write & feature flags** for a smooth transition, batch backfills, monitor aggressively, and have a clear rollback (snapshot/promote) plan. For RDS, prefer gh-ost/replicas because of privilege limits and always validate on the specific engine/version.”

---

What does “**zero downtime**” mean during a schema migration?

- A. The migration happens instantly
- B. Users can continue using the application without noticing any outage
- C. The database does not take backups during migration
- D. Writes are paused but reads continue

✓ **Correct Answer: B.**

💡 **Explanation:** **Zero-downtime** migration means the app continues to function normally — no interruptions for end users during schema or data changes.

---

Which of the following statements is **TRUE** about safe schema changes?

- A. Always drop old columns first before adding new ones
- B. Avoid making backward-incompatible changes until all app instances are updated
- C. Always perform ALTER TABLE directly on production
- D. Add a NOT NULL column with default immediately

✓ **Correct Answer: B.**

💡 **Explanation:** Safe schema changes are **backward-compatible first** — add before removing, ensuring both old and new app versions work together.

---

On MySQL RDS, which tool can perform **online schema migrations** without locking tables?

- A. mysqldump
- B. gh-ost or pt-online-schema-change
- C. AWS DMS
- D. CloudFormation

✅ **Correct Answer: B.**

💡 **Explanation:** **gh-ost** and **pt-online-schema-change** are online migration tools that alter tables without full table locks — ideal for RDS.

---

What is the **safest first step** before running any schema migration on RDS production?

- A. Disable Multi-AZ
- B. Take a recent snapshot or backup
- C. Drop unused tables to free space
- D. Disable automatic backups

✅ **Correct Answer: B.**

💡 **Explanation:** A **snapshot** provides a rollback point if something goes wrong during schema modification.

---

During a zero-downtime migration, **dual-write logic** means:

- A. Writing the same data to two DBs simultaneously for failover
- B. Writing updates to both old and new schema columns temporarily
- C. Logging DB writes twice for redundancy
- D. Using replicas for read/write separation

✅ **Correct Answer: B.**

💡 **Explanation:** **Dual-write** ensures consistency between old and new schema fields during transition — critical for non-blocking migrations.

---