

@devopschallengehub



## How will you plan and implement rollback strategies for RDS ?

### 1) Principles first (what rollback must guarantee)

- **Minimal data loss** — prefer PITR or replica-based rollback to avoid losing recent writes.
- **Minimal downtime** — aim for quick DNS/app cutover or promotion of a replica.
- **Safe and tested** — every rollback path must be practiced in staging.
- **Clear owner & runbook** — who executes, who approves, who monitors.

Whenever I plan a database change or migration, I always think of rollback from the start. My goal is simple — if something goes wrong, I should be able to bring the system back to a good state quickly and safely.

First, before any schema change or deployment, I always take a **manual snapshot** of the RDS instance. This gives me a clean restore point.

I also make sure **automated backups** and **point-in-time recovery** are enabled, so I can restore the DB to any exact time if needed.

Now, if an issue happens after deployment, I choose the rollback method based on what broke:

1. **If it's just an app issue** (like a bug in the new release) — I simply roll back the app version or turn off the feature flag. No DB change is needed.
2. **If the app and DB are out of sync** (maybe the new app expects a new column) — I roll back the app and point it back to the previous DB schema. If the old schema isn't compatible, I restore from the snapshot I took before migration.
3. **If data got corrupted or deleted by mistake** — I use **Point-in-Time Restore (PITR)** to create a new RDS instance from a few minutes before the issue. Then I point the app to this new instance.
4. **If I have a read replica ready** — I can **promote the replica** to be the new primary. This is faster and usually causes less data loss.
5. **For big migrations** — I prefer **blue-green deployment**. That means I create a new RDS (blue) with the new schema, test it, and switch traffic. If something fails, rollback is easy — I just switch back to the old (green) DB.

After rollback, I always **test the key functions** — can users log in, can we fetch data, are reports working, etc. Then I monitor the logs and CloudWatch metrics for a few hours.

Finally, we do a short review: what went wrong, how can we improve the rollback plan next time, and update our runbook.

So in short —

- 👉 I prepare before making changes (snapshot + backups).
- 👉 I pick the right rollback method based on the problem.
- 👉 I test, monitor, and learn from every rollback.”

---

What is the **primary goal** of a rollback plan for RDS?

- A. To rebuild the RDS instance from scratch
- B. To restore service quickly with minimal data loss and downtime
- C. To delete corrupted databases immediately
- D. To disable all automated backups

✅ **Correct Answer: B.**

💡 **Explanation:** A rollback plan ensures **fast recovery** while maintaining **data integrity** and **availability** after a failure or bad deployment.

---

Before making a schema change or major RDS update, what should you **always do first**?

- A. Disable backups to improve performance
- B. Take a manual snapshot of the RDS instance
- C. Promote a replica to primary
- D. Modify the DB parameter group

✅ **Correct Answer: B.**

💡 **Explanation:** Always take a **manual snapshot** before risky operations to have a **safe restore point** for rollback.

---

What does a **Point-In-Time Restore (PITR)** do?

- A. Rolls back schema changes without data loss
- B. Creates a new RDS instance from backups and logs to a specific timestamp
- C. Restores the same DB in place
- D. Promotes a replica automatically

✅ **Correct Answer: B.**

💡 **Explanation:** PITR restores data **to an exact second** before failure by replaying backup logs — creating a **new DB instance**.

---

If an application release introduces a bug but the DB schema is unchanged, what's the **fastest rollback path**?

- A. PITR to before the release
- B. Restore from snapshot
- C. Redeploy the previous application version or disable the feature flag
- D. Reboot the RDS instance

✅ **Correct Answer: C.**

💡 **Explanation:** When DB schema is still compatible, **application rollback** via **previous version** or **feature flag** is fastest — no DB rollback needed.

---

Which rollback method allows you to **minimize data loss** by switching to an almost-synced replica?

- A. PITR restore
- B. Blue/Green deployment
- C. Promote a read replica
- D. Manual snapshot restore

✅ **Correct Answer: C.**

💡 **Explanation:** Promoting a read replica provides a **fast failover** with minimal data loss (depending on **replication lag**).

---

What is the **main limitation** of restoring from a manual snapshot compared to PITR?

- A. It restores slower
- B. It can only restore to the exact snapshot time, not a specific second
- C. It requires deletion of the existing DB first
- D. It cannot be automated

✅ **Correct Answer: B.**

💡 **Explanation:** A **manual snapshot** restores DB state **exactly** as captured — PITR offers **time-based precision**.

---

What's the **main advantage** of a Blue/Green (clone + cutover) RDS deployment for rollback?

- A. No need for backups
- B. Rollback is instant — just switch traffic back to the old environment
- C. Automatic data synchronization without replication lag
- D. Reduces storage cost

✅ **Correct Answer: B.**

💡 **Explanation:** In **Blue/Green**, rollback is a **simple DNS or connection switch** — fast, low risk, ideal for major upgrades.

---