

@devopschallengehub



## How would you debug RDS connectivity issues or performance bottlenecks?

When I see an issue with RDS, I start with the basics — can the application even reach the database?

First, I check **connectivity**.

From an EC2 or bastion host, I'll try to connect to the DB using its endpoint and port — for example, `telnet <db-endpoint> 3306`.

If that fails, I look at **security groups**, **VPC routes**, and **subnet settings** to make sure traffic is allowed.

Sometimes it's just a **DNS** issue or the app is using an old cached DNS entry — restarting the app often fixes that.

If we use **IAM authentication**, I also check if the token expired. (Normally, to connect to a database (like MySQL or PostgreSQL), you use a **username** and **password** that are stored **inside the database** itself.

To solve this, **AWS allows you to use IAM authentication**. This means — instead of a fixed DB password, your app uses a **temporary token** issued by **AWS IAM** (Identity and Access Management).)

Once I confirm the app can connect, I move to **performance**.

Here I rely on **CloudWatch metrics** first — CPU usage, storage space, number of connections, and IOPS.

If CPU is very high or there's a big spike in IOPS, I open **Performance Insights** to see which queries are consuming the load.

This quickly shows if one or two queries are slowing everything down.

If queries are slow, I check the **slow query logs** or run **EXPLAIN** on the query to see if it's missing an index or doing a full table scan.

Sometimes the issue is not the query but too many open connections — so I also check the connection pool on the application side.

If it's a short-term emergency, I might take quick actions like:

- Killing a runaway query.
- Scaling the DB instance to a larger size.
- Increasing IOPS or routing some read traffic to a read replica.

After stabilizing things, I focus on **long-term fixes** — like optimizing queries, adding the right indexes, using caching (Redis), and setting up proper connection pooling.

I also tune DB parameters like buffer size or max connections and keep monitoring through CloudWatch and Performance Insights.

In short —

- 👉 I start with **connectivity** (network, SGs, DNS).
- 👉 Then I look at **metrics and logs** to find the bottleneck (CPU, IO, locks, or queries).
- 👉 I apply **quick fixes** to stabilize, and then **optimize** for long-term performance.”

---

**Q:** What is the *first* thing you should check when an application cannot connect to an RDS instance?

- A. Performance Insights for slow queries
- B. Security Groups, NACLs, and port (e.g., 3306/5432)
- C. Increase instance size
- D. Promote a read replica

✅ **Correct: B**

**Why:** Connectivity failures are usually network/security issues — verify SG rules, NACLs, and the DB port before anything else.

---

**Q:** Which CloudWatch metric tells you the remaining disk capacity of your RDS instance?

- A. CPUUtilization
- B. DatabaseConnections
- C. FreeStorageSpace
- D. ReadIOPS

✅ **Correct: C**

**Why:** FreeStorageSpace shows available storage — alert when it’s low to avoid DB stalls.

---

**Q:** Which simple command from a bastion host tests basic network reachability to a MySQL RDS endpoint?

- A. ps -aux | grep mysql
- B. telnet mydb.xxxx.rds.amazonaws.com 3306
- C. aws rds describe-db-instances
- D. SELECT 1;

✅ **Correct: B**

**Why:** telnet (or nc -vz) checks whether the TCP port is reachable from the network location.

---

**Q:** What does a high ReplicaLag metric indicate?

- A. The primary instance is down
- B. Read replica is falling behind writes — reads may be stale
- C. CloudWatch is delayed
- D. The DB has too many connections

✅ **Correct: B**

**Why:** ReplicaLag in seconds shows replication delay; high values mean stale read replicas.

---

**Q: If CPU Utilization is low but queries are slow, what is the most likely bottleneck?**

- A. CPU-bound queries
- B. I/O (IOPS)/disk latency or connection saturation
- C. Lack of backups
- D. DNS TTL issues

✅ **Correct: B**

**Why:** Low CPU with slow queries points to I/O bottlenecks, lock contention, or connection issues rather than compute.

**I/O Bottleneck — Input/Output bottleneck**

**I/O** means **Input/Output**, or simply — **reading and writing data** to disk.

When you run a query, the database has to:

- Read data from disk (storage)
- Or write data (like when inserting or updating rows)

If the disk is **slow**, or the DB is doing **too many reads/writes** at once, queries start waiting.

This is called an **I/O bottleneck** — the storage can't keep up.

Databases use **locks** to keep data safe when multiple users try to change it at the same time.

For example:

- Query A is updating a row in the “orders” table.
- Query B tries to update the same row at the same time.
- Query B must **wait** until Query A finishes.

If this happens too often, many queries are **waiting for locks** — this is called **lock contention**.

---

**Q: Which RDS feature gives **query-level** visibility (top SQL, waits) to identify slow queries?**

- A. CloudTrail
- B. Enhanced Monitoring
- C. Performance Insights
- D. DB Subnet Group

✅ **Correct: C**

**Why:** Performance Insights shows top SQL by load, wait events, and helps find problematic queries.

---