



Describe your experience with RDS in a production environment ?

In production, I've used Amazon RDS — mainly MySQL and PostgreSQL — as the main database for our applications. I usually provision it using **Terraform**, so the database setup is fully automated and version-controlled. RDS sits inside a **private VPC** with proper **security groups**, and for all production systems, we enable **Multi-AZ** for high availability.

In our **CI/CD pipeline**, the database part is handled carefully.

The RDS instance itself is not recreated in every build — it's managed separately as infrastructure-as-code.

When developers add a new feature, sometimes they also need to **change the database** — for example, add a new column like `phone_number` in the `users` table.

Instead of doing that manually in the RDS console or with SQL commands, we use a **migration tool** such as **Flyway** or **Liquibase**.

So whenever we deploy a new release, the pipeline first runs migrations in **staging**, runs automated tests, and only then applies them to **production**.

If a migration might be risky, we take a **blue/green approach** — clone the DB, apply changes, test, and then cut over to the new one.

For **high availability**, all production RDS databases run in **Multi-AZ mode**.

That means AWS automatically maintains a standby database in another availability zone with synchronous replication.

If the main DB goes down, failover happens automatically — usually within 30–60 seconds.

We've tested this using the **"Reboot with failover"** option during off-peak hours to make sure the app reconnects properly.

We also added **connection retry logic** and **pooling (PgBouncer or HikariCP)** in the app so users don't notice small failovers.

For **disaster recovery**, we set up **cross-region read replicas** — for example, our main DB was in Mumbai (ap-south-1) and a replica was in Virginia (us-east-1).

If the primary region ever fails, we can **promote the replica** and point the app to the new endpoint.

We also copy automated **snapshots** every night to another region as a backup.

Our DR plan is documented — it includes steps like promoting the replica, updating DNS, and running smoke tests to confirm recovery.

We even do **quarterly DR drills** to practice this and make sure our RTO and RPO targets are met.

So in short —

I manage RDS as **infrastructure-as-code**, integrate it safely into the **CI/CD pipeline** for schema changes, ensure **high availability** with Multi-AZ and failover testing, and maintain **disaster recovery** through cross-region replicas and snapshot backups.

That setup gives us both uptime and confidence that we can recover from failures quickly.

Which RDS configuration provides **automatic failover** to a standby instance in another Availability Zone?

- A. Read Replica
- B. Multi-AZ Deployment
- C. Aurora Serverless
- D. Parameter Group

✅ **Correct Answer: B. Multi-AZ Deployment**

💡 **Explanation:** Multi-AZ provides **synchronous replication** to a standby DB in another AZ. On failure, RDS automatically promotes the standby.

During an RDS failover, what happens to the database endpoint?

- A. The endpoint IP changes, but the DNS name remains the same.
- B. The endpoint DNS name changes permanently.
- C. You must manually update application config.
- D. Failover requires manual endpoint remapping.

✅ **Correct Answer: A.**

💡 **Explanation:** RDS updates the **DNS record** to point to the new primary. The DNS name stays constant, ensuring app transparency.

Which RDS feature should you use to replicate data across AWS **regions** for disaster recovery?

- A. Multi-AZ deployment
- B. Read Replica (Cross-Region)
- C. Elastic Beanstalk
- D. Parameter Group replication

✅ **Correct Answer: B. Cross-Region Read Replica**

💡 **Explanation:** Multi-AZ is intra-region; **Cross-region read replicas** enable asynchronous replication for regional DR setups.

How do you typically integrate RDS schema migrations into a CI/CD pipeline?

- A. Manually apply scripts after deployment
- B. Include migration tools like Flyway/Liquibase as a pipeline stage
- C. Use AWS Lambda triggers for DB updates
- D. Apply migrations via AWS Config

✅ **Correct Answer: B.**

💡 **Explanation:** Tools like **Flyway, Liquibase, or Alembic** are commonly integrated into CI/CD to apply schema changes automatically and safely.

Which Infrastructure-as-Code tool is most commonly used to **provision RDS** in automated pipelines?

- A. Chef
- B. Terraform
- C. Puppet
- D. AWS Inspector

✅ **Correct Answer: B. Terraform**

💡 **Explanation:** Terraform is widely used for managing RDS configurations — VPC, subnet group, security, Multi-AZ, backups — through IaC.

What is the **primary difference** between Multi-AZ and Read Replica in RDS?

- A. Multi-AZ improves read scaling; Read Replica improves availability.
- B. Multi-AZ is synchronous for HA; Read Replica is asynchronous for read scaling.
- C. Both are synchronous replication methods.
- D. Multi-AZ is for analytics; Read Replica is for DR.

✅ **Correct Answer: B.**

💡 **Explanation:** Multi-AZ → **synchronous**, for HA; Read Replica → **asynchronous**, for **read scaling or cross-region DR**.

In a production-grade CI/CD pipeline, how are **risky schema changes** (like column drops or type changes) typically handled?

- A. Directly deployed to production.
- B. Using blue/green or clone → migrate → cutover strategy.
- C. Executed via Lambda.
- D. Ignored during deployment.

✅ **Correct Answer: B.**

💡 **Explanation:** A **blue/green migration** isolates the schema change to a cloned DB, validated before final cutover — ensuring zero downtime and rollback safety.
